
SPOTON:

A BATCH COMPUTING SERVICE FOR THE SPOT MARKET

**Supreeth Subramanya, Tian Guo, Prateek Sharma,
David Irwin, Prashant Shenoy**

University of Massachusetts Amherst

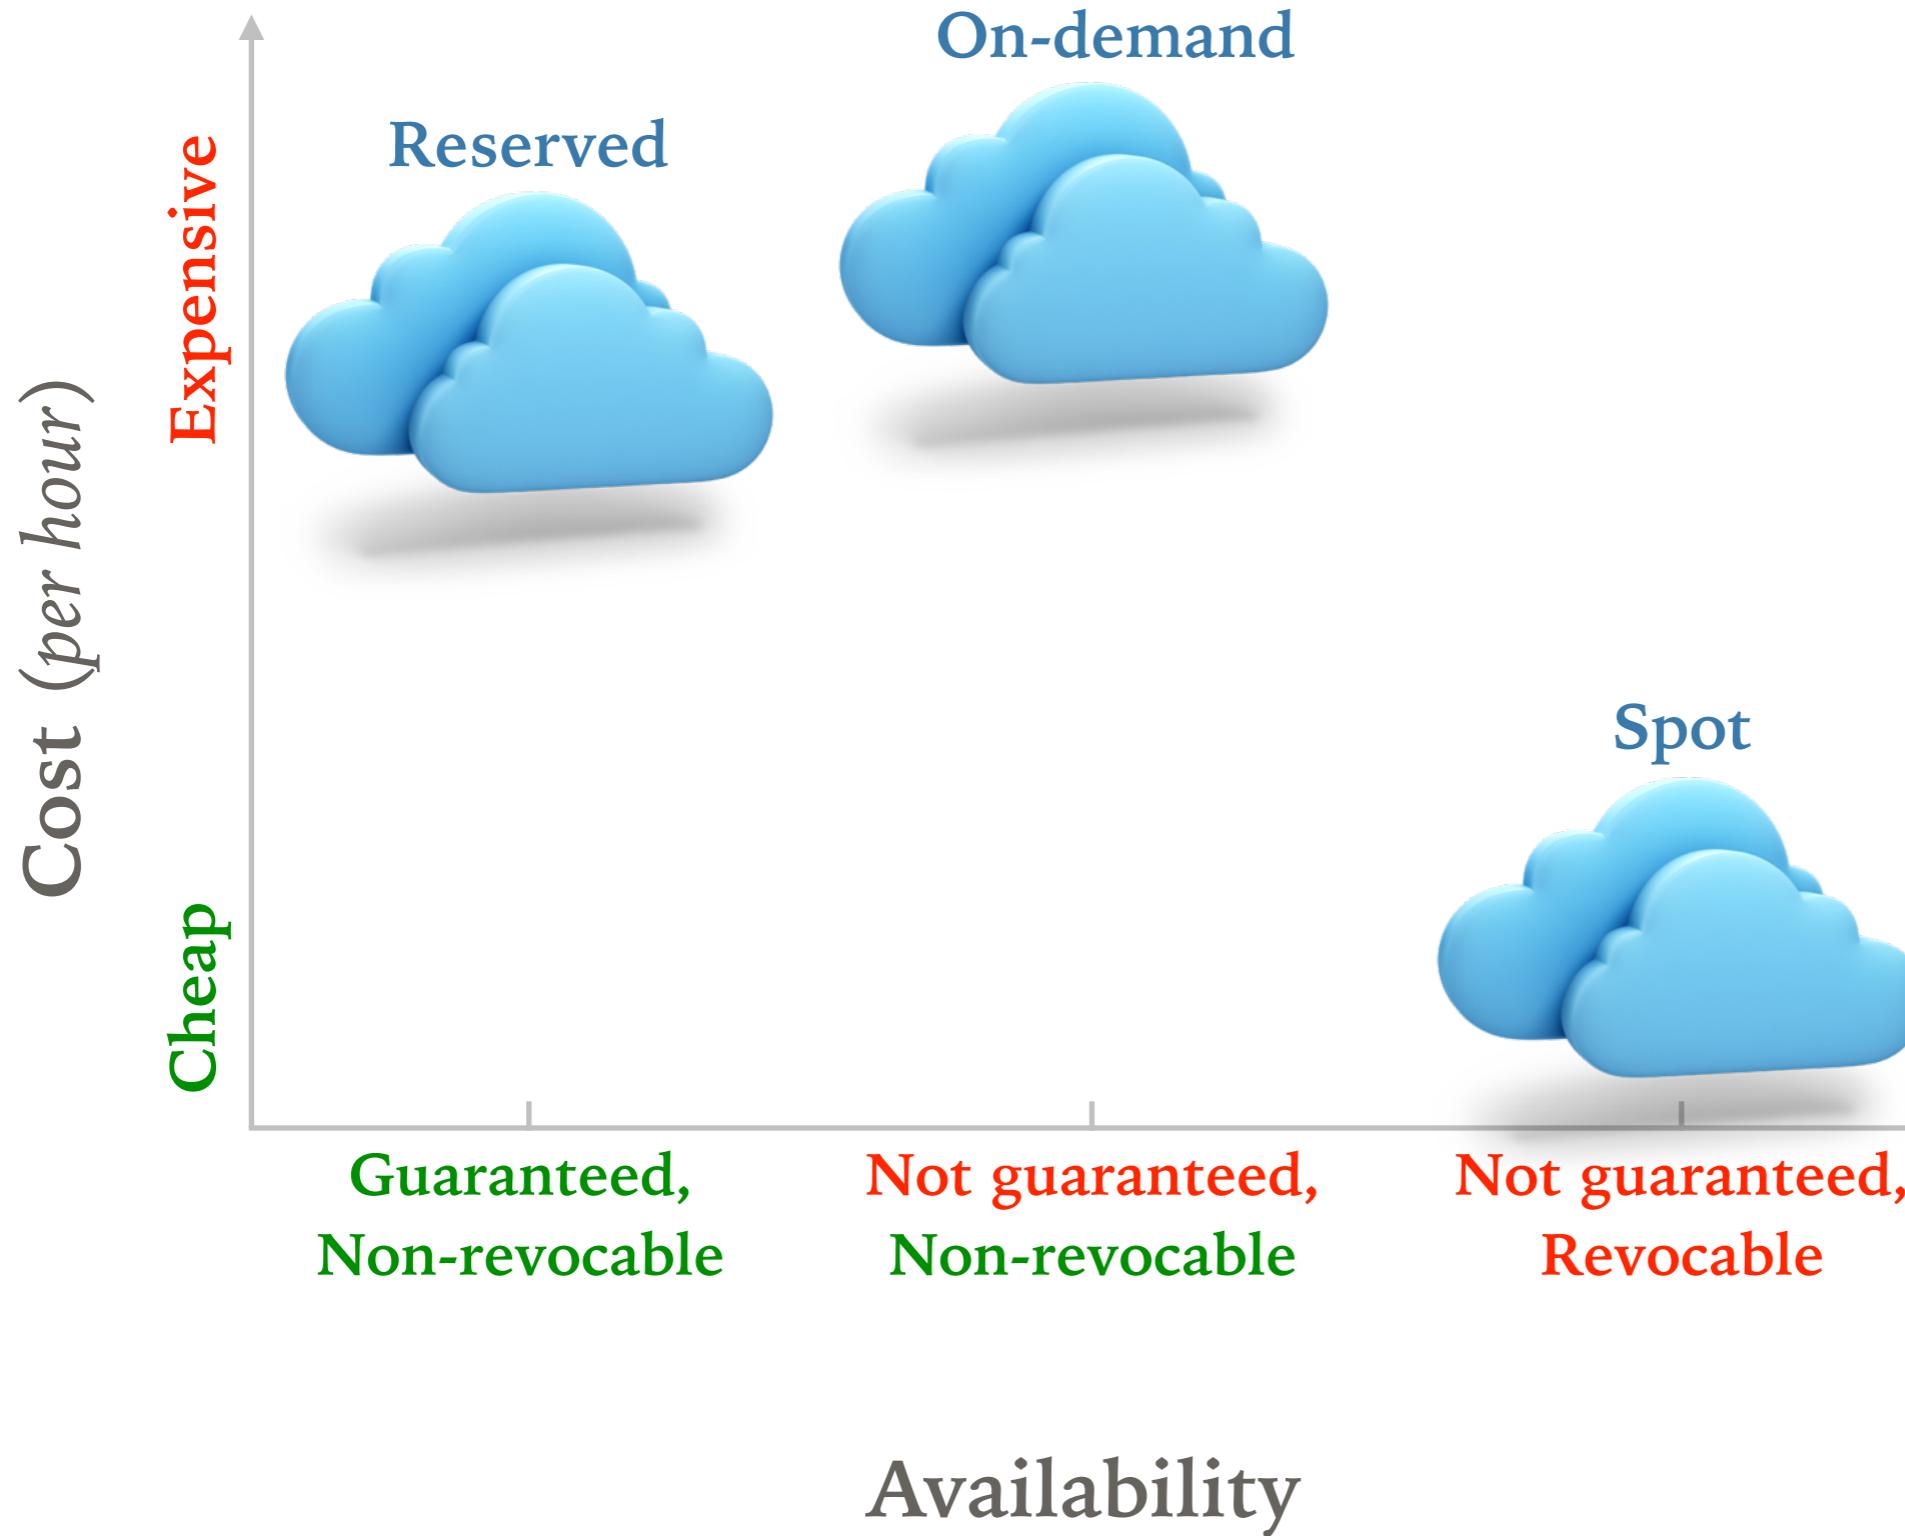
INFRASTRUCTURE CLOUD

On-demand



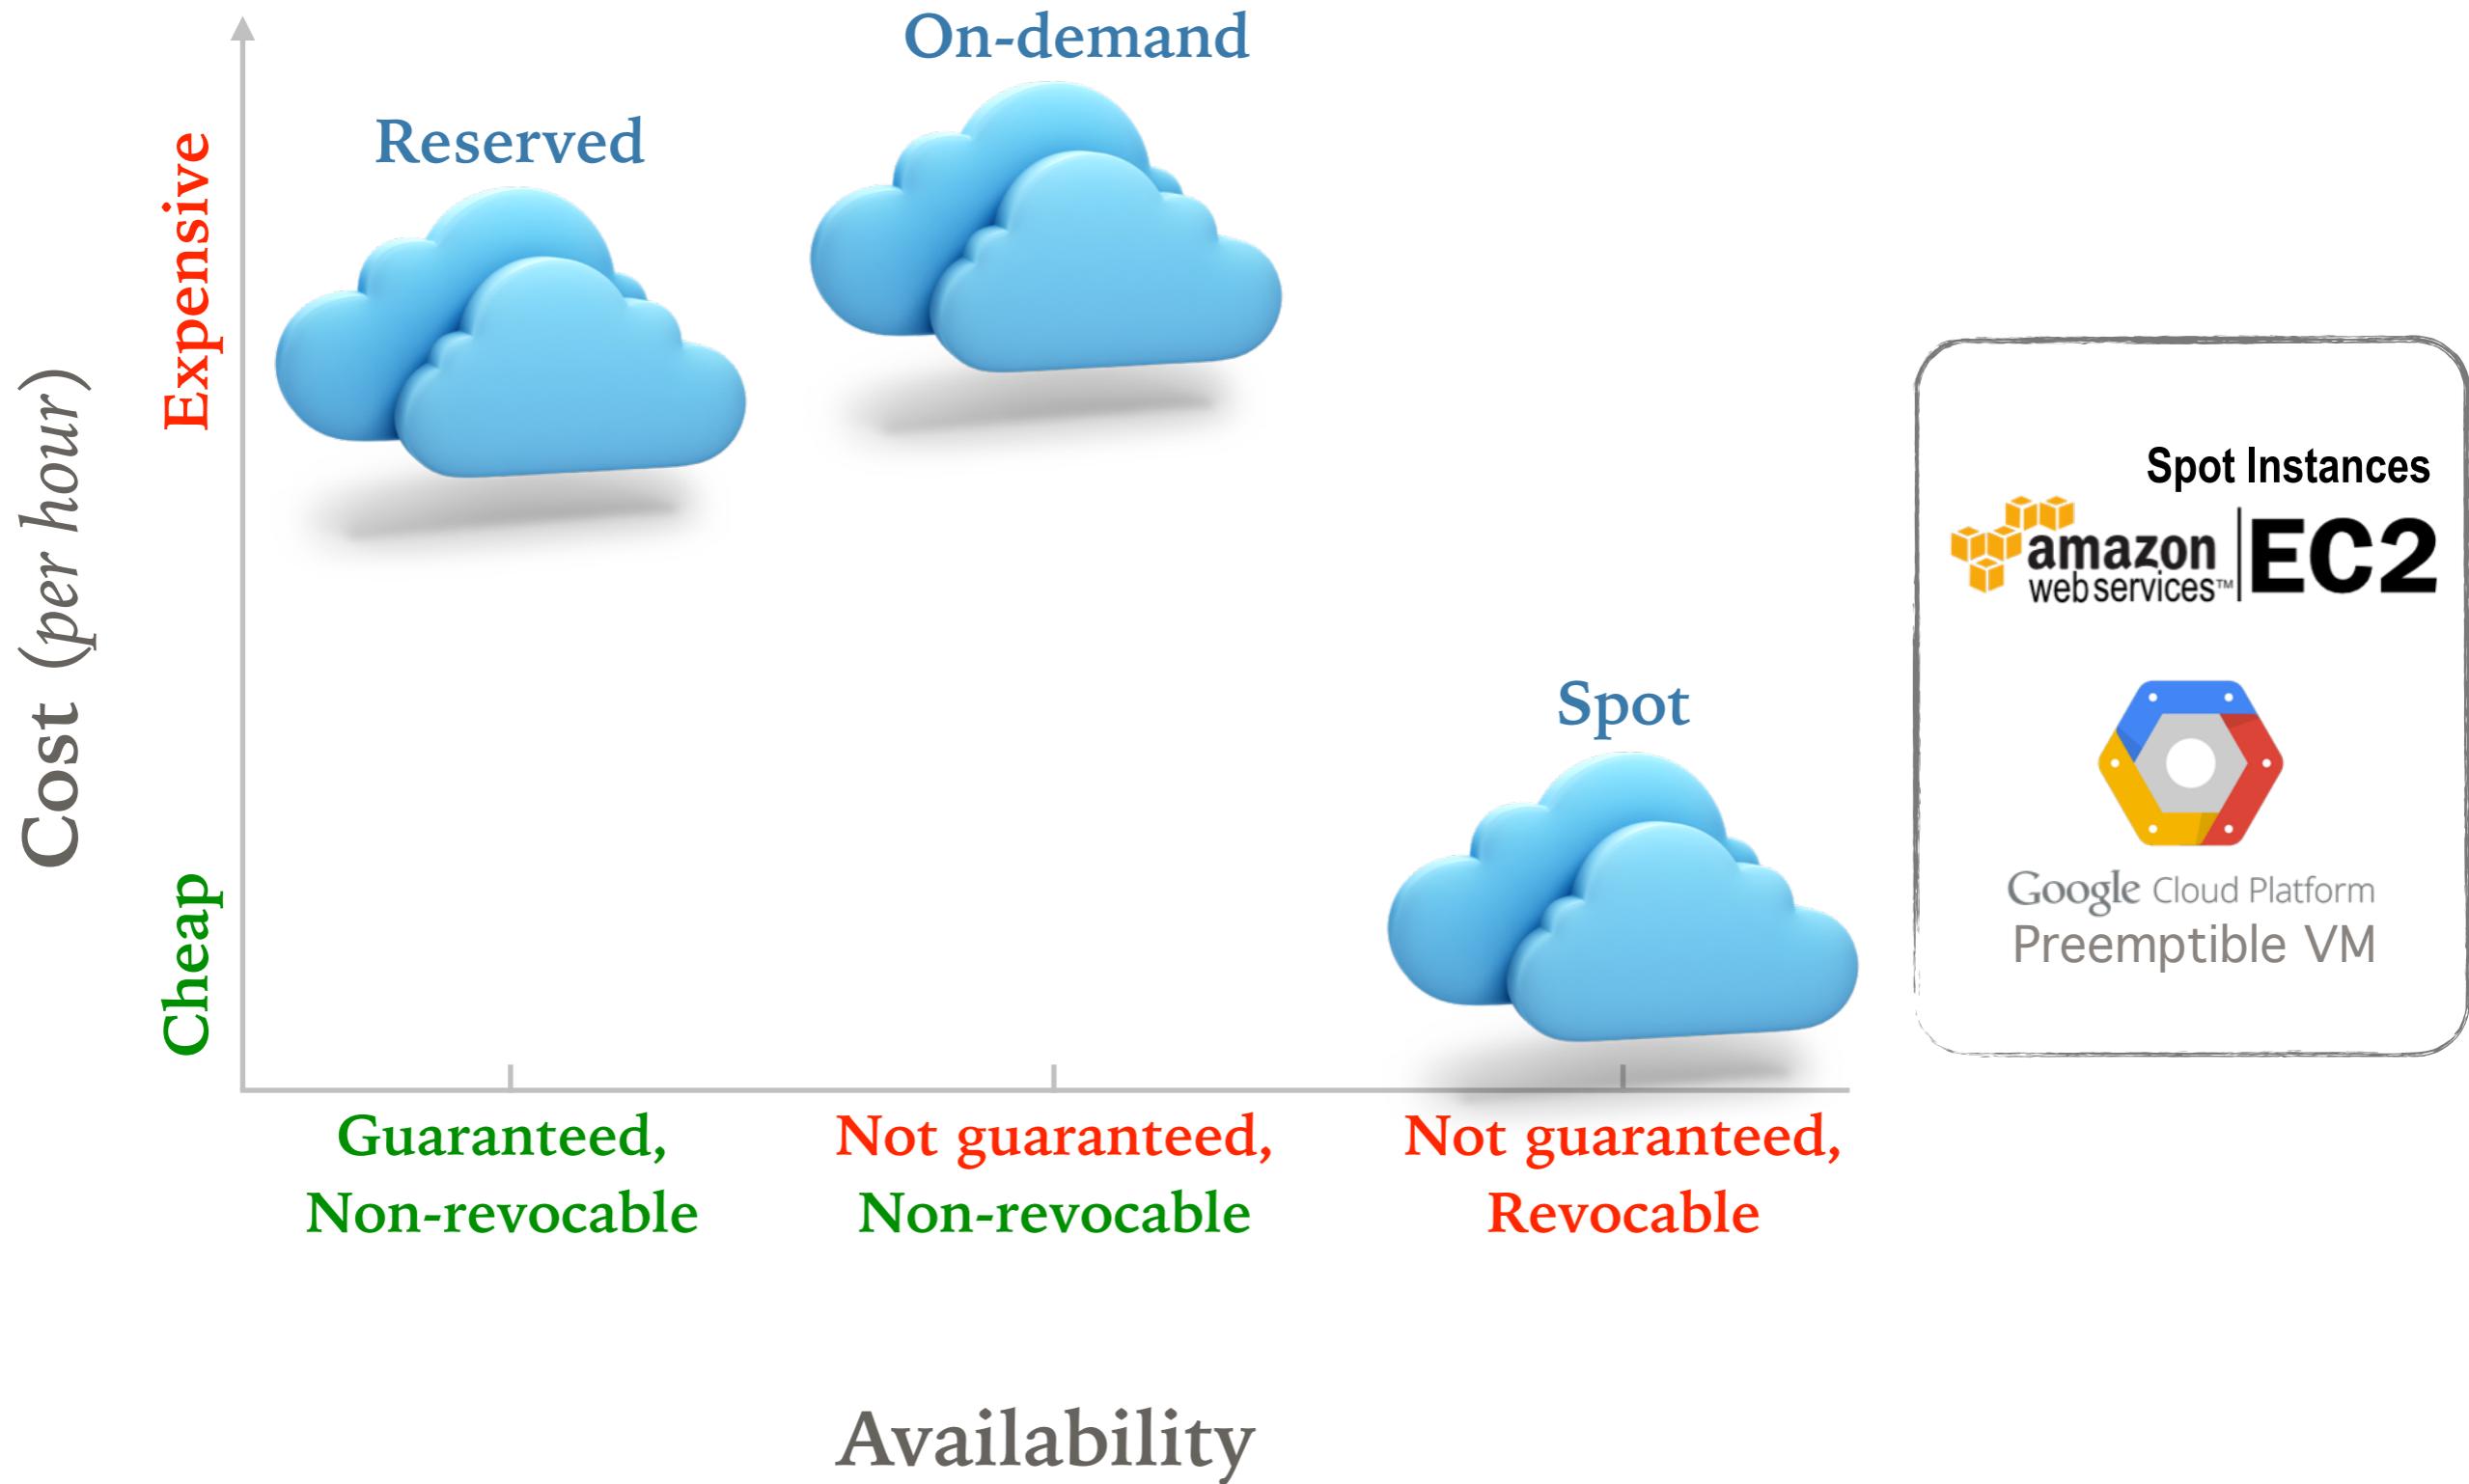
INFRASTRUCTURE CLOUD

Cost vs. Availability Tradeoff in IaaS Clouds



INFRASTRUCTURE CLOUD

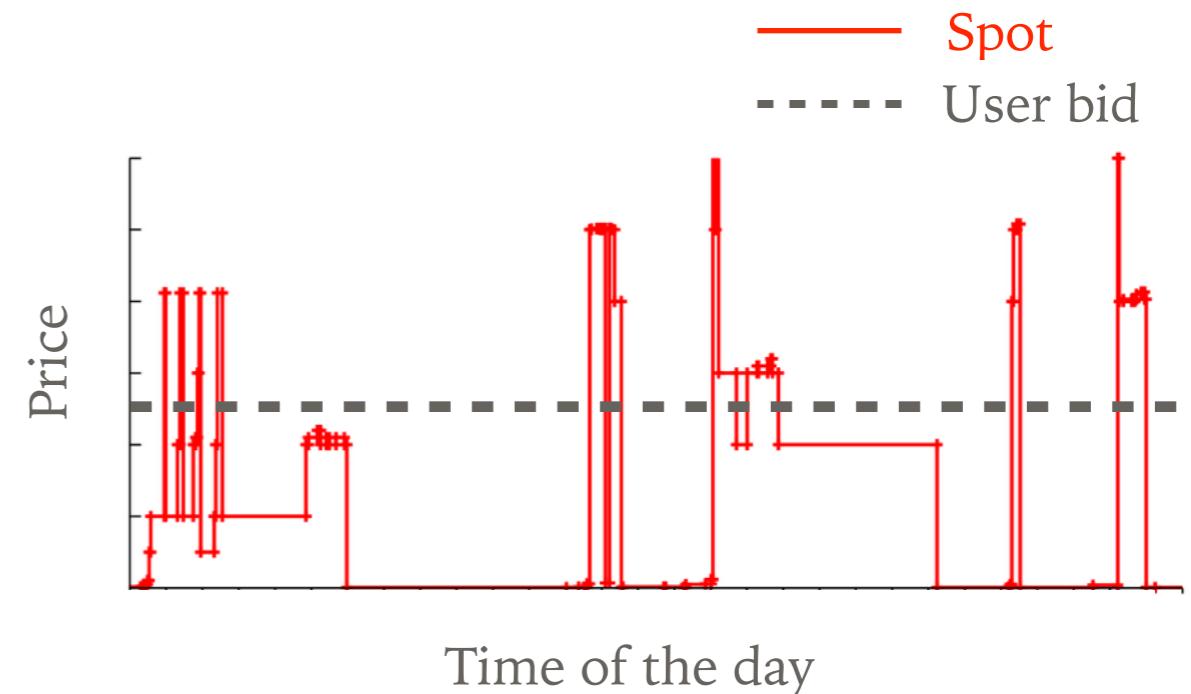
Cost vs. Availability Tradeoff in IaaS Clouds



SPOT MARKETS

Amazon EC2

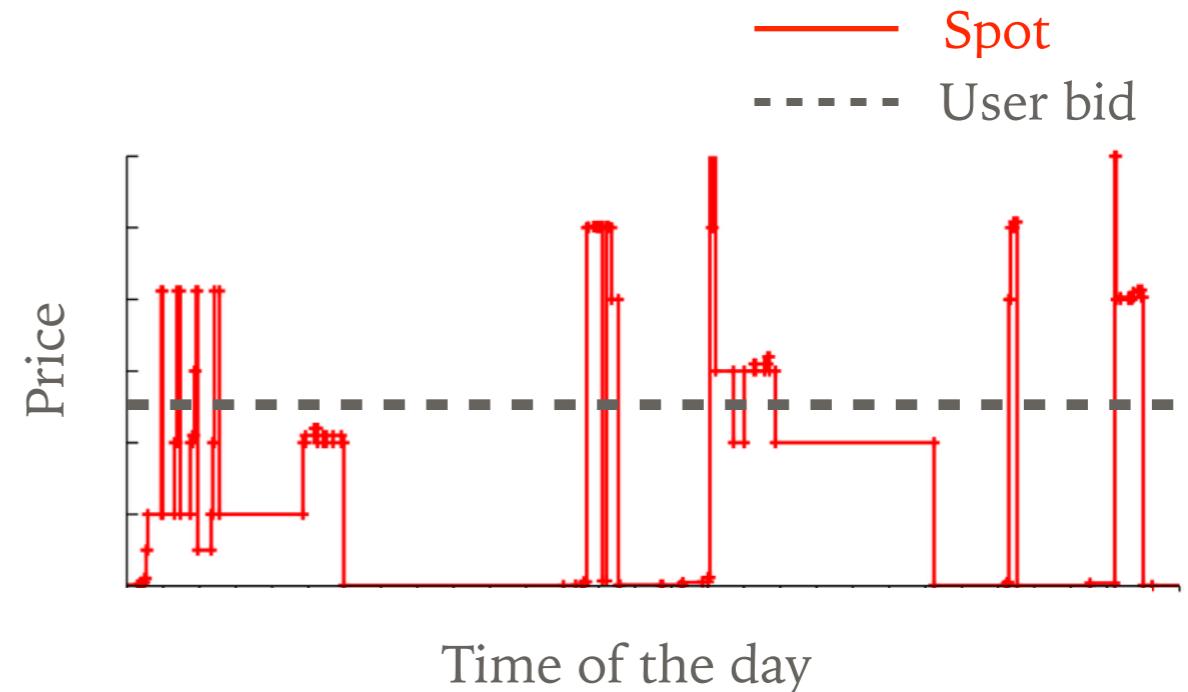
- ▶ **Bid** in a 2nd price auction
- ▶ **Acquire** when bid > spot price
- ▶ **Terminate** when spot price > bid



SPOT MARKETS

Amazon EC2

- ▶ **Bid** in a 2nd price auction
- ▶ **Acquire** when bid > spot price
- ▶ **Terminate** when spot price > bid



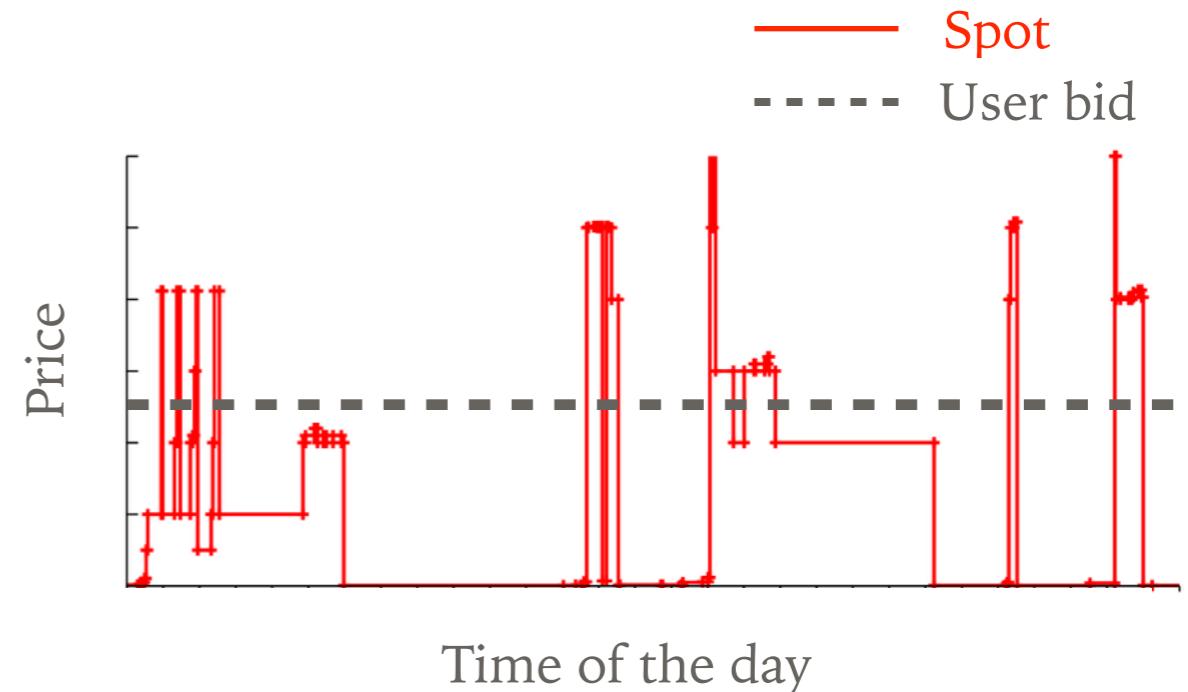
How do we mitigate the impact of revocation?

1. Raise the bid
2. Employ fault-tolerance mechanisms

SPOT MARKETS

Amazon EC2

- ▶ **Bid** in a 2nd price auction
- ▶ **Acquire** when bid > spot price
- ▶ **Terminate** when spot price > bid



How do we mitigate the impact of revocation?

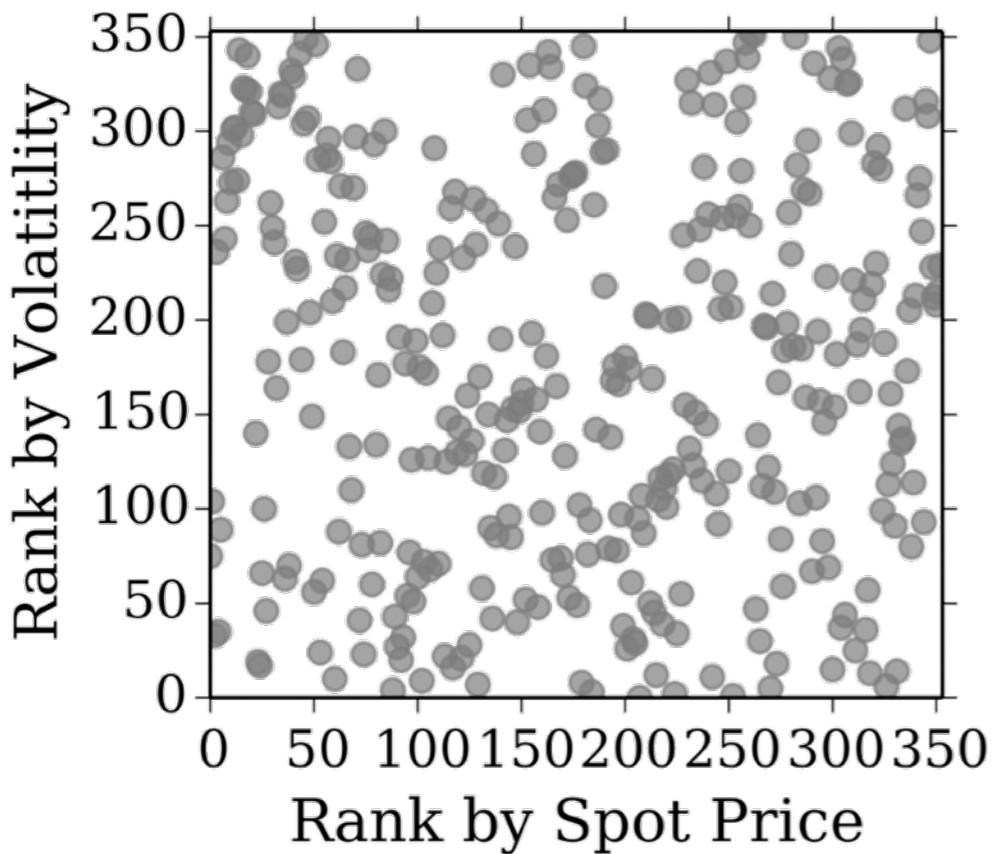
1. Raise the bid
2. Employ fault-tolerance mechanisms

CHALLENGES – SPOT MARKET COMPLEXITY

Amazon EC2 operates ~**4000** spot markets

CHALLENGES – SPOT MARKET COMPLEXITY

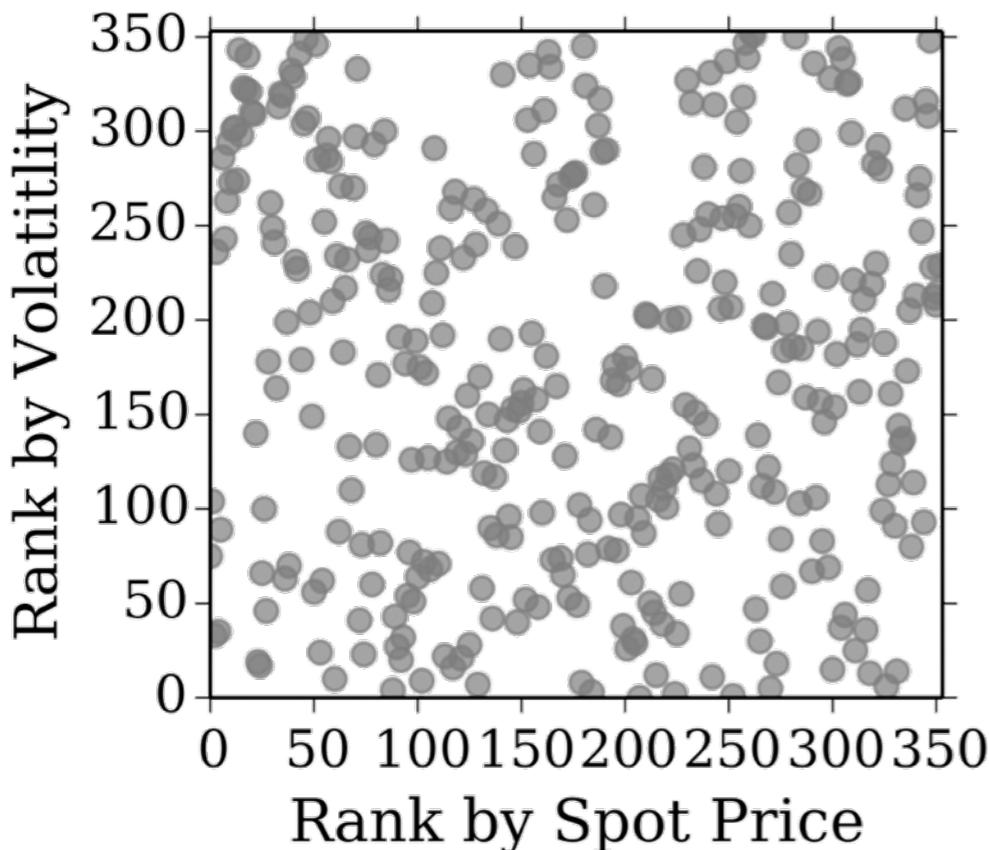
Amazon EC2 operates ~ 4000 spot markets



Scatterplot of ranks for EC2 spot markets

CHALLENGES – SPOT MARKET COMPLEXITY

Amazon EC2 operates ~ 4000 spot markets



Scatterplot of ranks for EC2 spot markets

Selecting an instance that yields
lowest cost per unit of computation
while also considering the
probability of revocation is complex

CHALLENGES – APPLICATION COMPLEXITY



CPU : IO	Working Set	Disk Type	Running Time
1:1	8GB	Remote	1 hour

Resource Vector

CHALLENGES – APPLICATION COMPLEXITY



CPU : IO	Working Set	Disk Type	Running Time
1:1	8GB	Remote	1 hour

Resource Vector



Cost	Revocation Rate	Fault-tolerance Mechanism
20% On-demand	2.4 per day	Checkpoint (every 900s)

Spot VM

CHALLENGES – APPLICATION COMPLEXITY



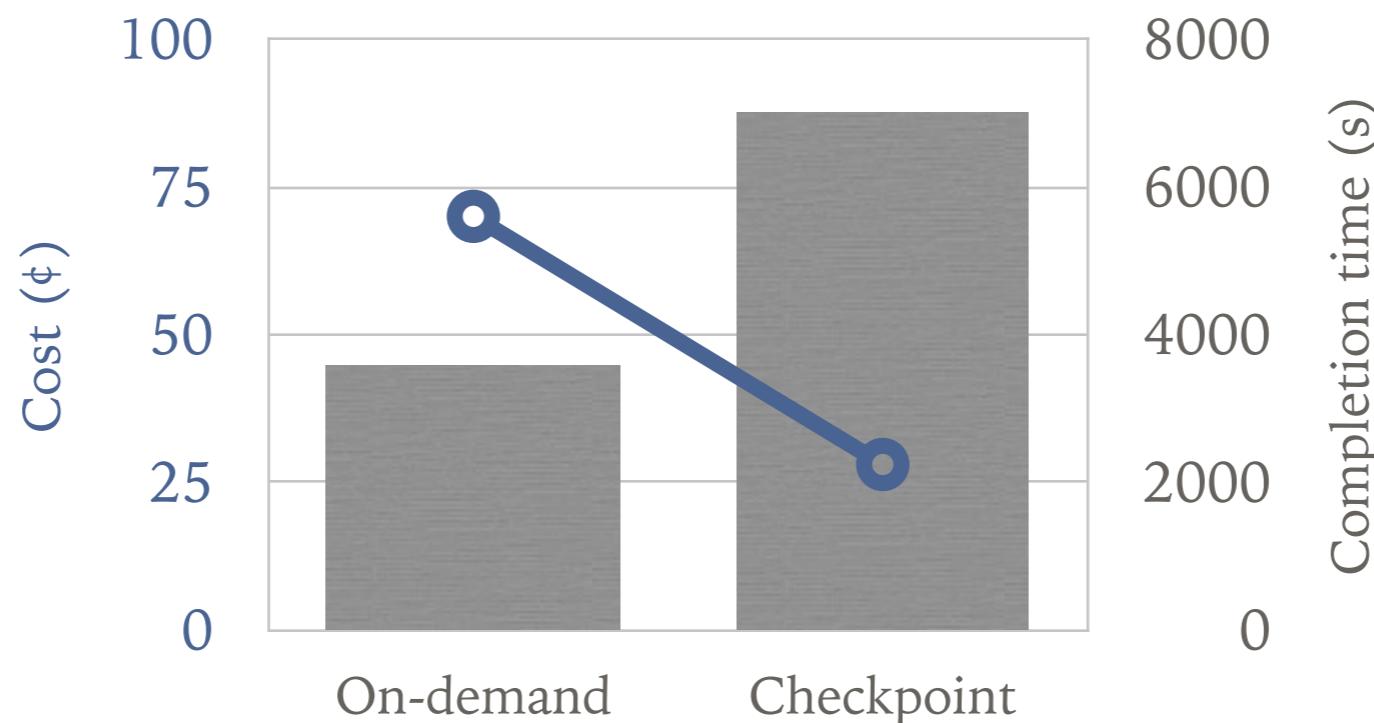
CPU : IO	Working Set	Disk Type	Running Time
1:1	8GB	Remote	1 hour

Resource Vector



Cost	Revocation Rate	Fault-tolerance Mechanism
20% On-demand	2.4 per day	Checkpoint (every 900s)

Spot VM



CHALLENGES – APPLICATION COMPLEXITY



Resource Vector

CPU : IO	Working Set	Disk Type	Running Time
1:1	8GB	Local	1 hour



Spot VM

Cost	Revocation Rate	Fault-tolerant mechanism
20% On-demand	2.4 per day	Replicate (deg=2)

CHALLENGES – APPLICATION COMPLEXITY



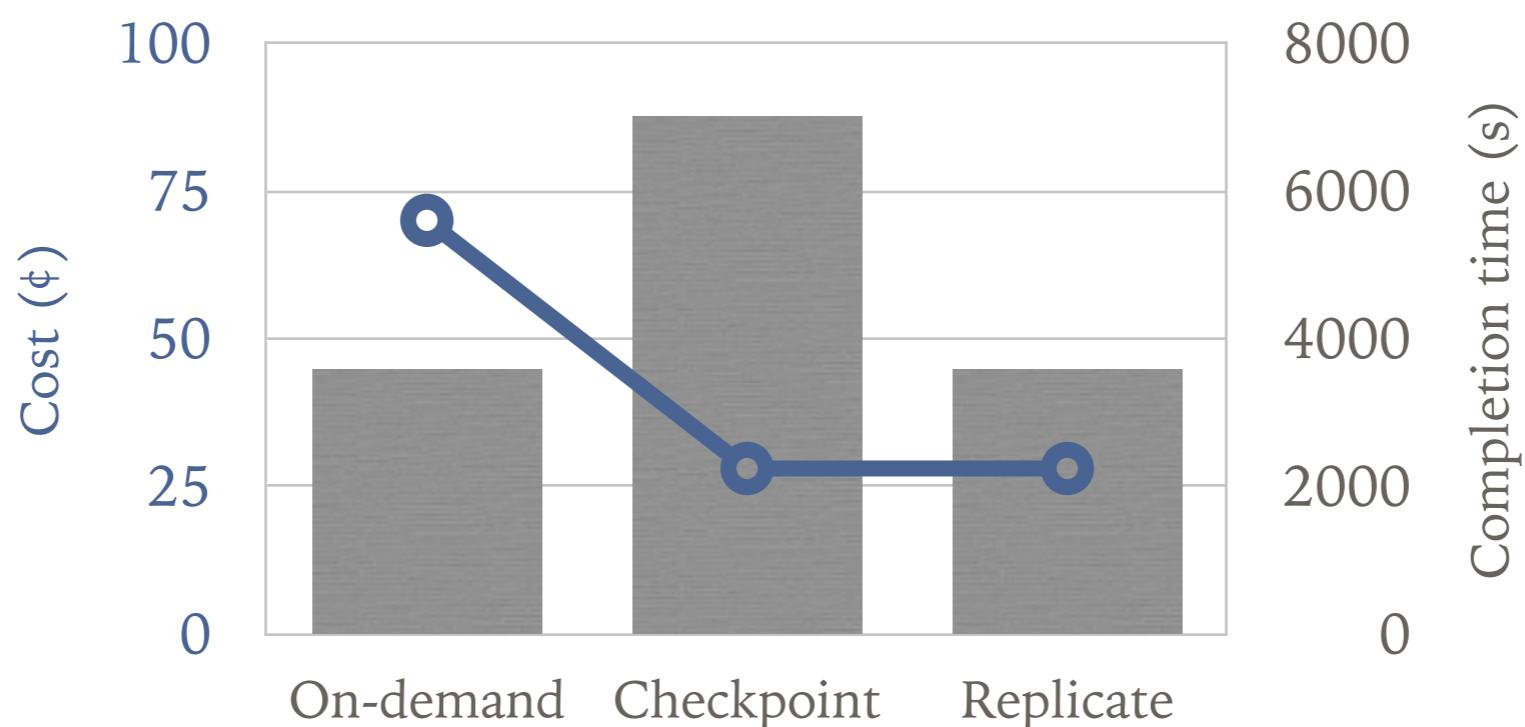
Resource Vector

CPU : IO	Working Set	Disk Type	Running Time
1:1	8GB	Local	1 hour



Spot VM

Cost	Revocation Rate	Fault-tolerance mechanism
20% On-demand	2.4 per day	Checkpoints (every 500s) Replicate (deg=2)



CHALLENGES – APPLICATION COMPLEXITY



CPU : IO	Working Set	Disk Type	Running Time
1:1	8GB	Local	1 hour

Resource Vector



Cost	Revocation Rate	Fault-tolerance Mechanism
20% On-demand	>24 per day	Checkpoint (every 900s)

Spot VM

CHALLENGES – APPLICATION COMPLEXITY



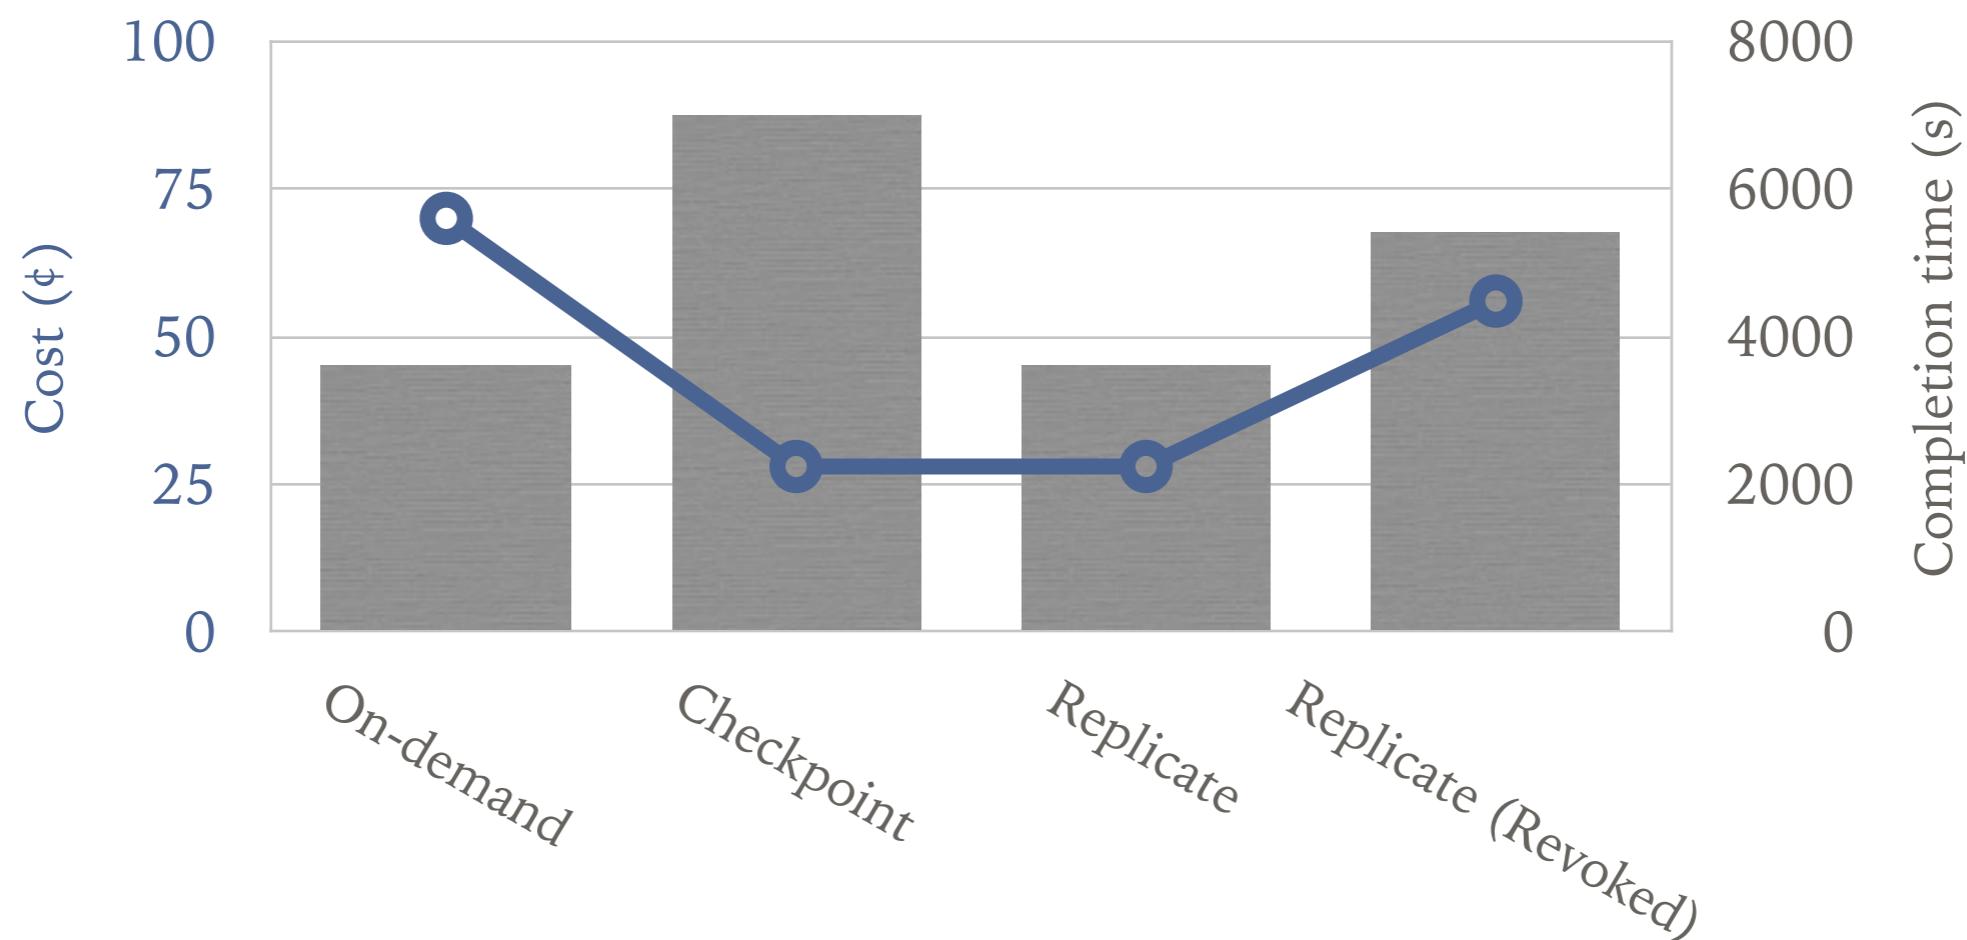
CPU : IO	Working Set	Disk Type	Running Time
1:1	8GB	Local	1 hour

Resource Vector



Cost	Revocation Rate	Fault-tolerance Mechanism
20% On-demand	>24 per day	Checkpoint (every 900s)

Spot VM



SPOTON: A BATCH COMPUTING SERVICE



SPOTON: A BATCH COMPUTING SERVICE



Service that accepts batch jobs from users and runs them on spot instances

Manages application and spot market complexity transparently

SPOTON: A BATCH COMPUTING SERVICE



Service that accepts batch jobs from users and runs them on spot instances

Manages application and spot market complexity transparently

Run batch jobs at *on-demand performance*
but paying *spot market price*

GREEDY SELECTION ALGORITHM

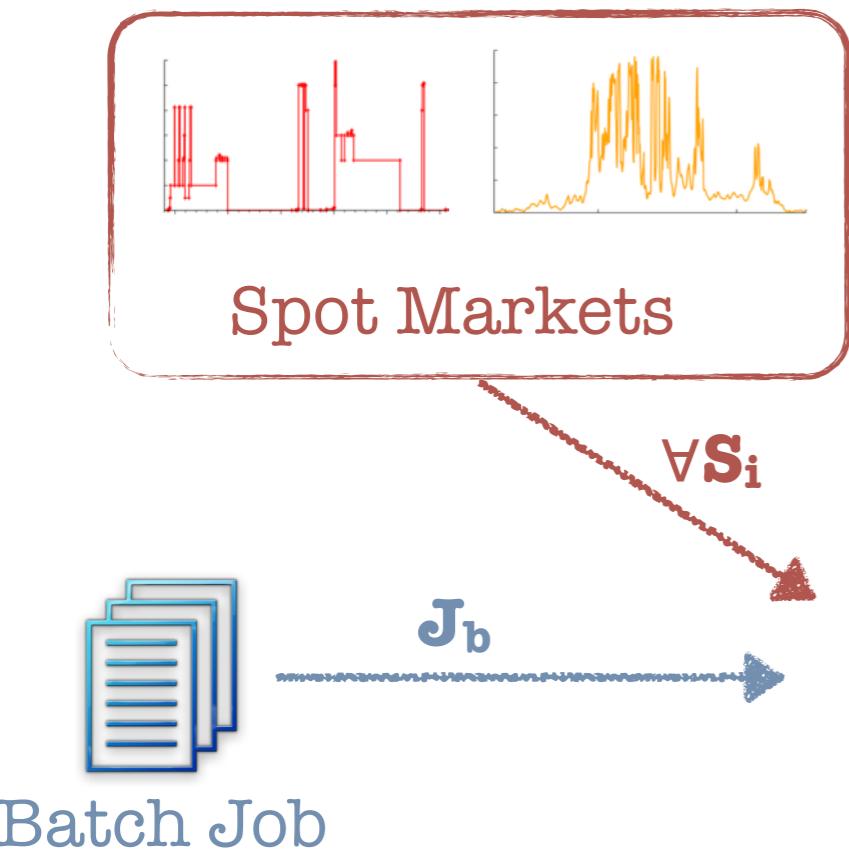
Selecting the best spot market and fault-tolerance mechanism



Batch Job

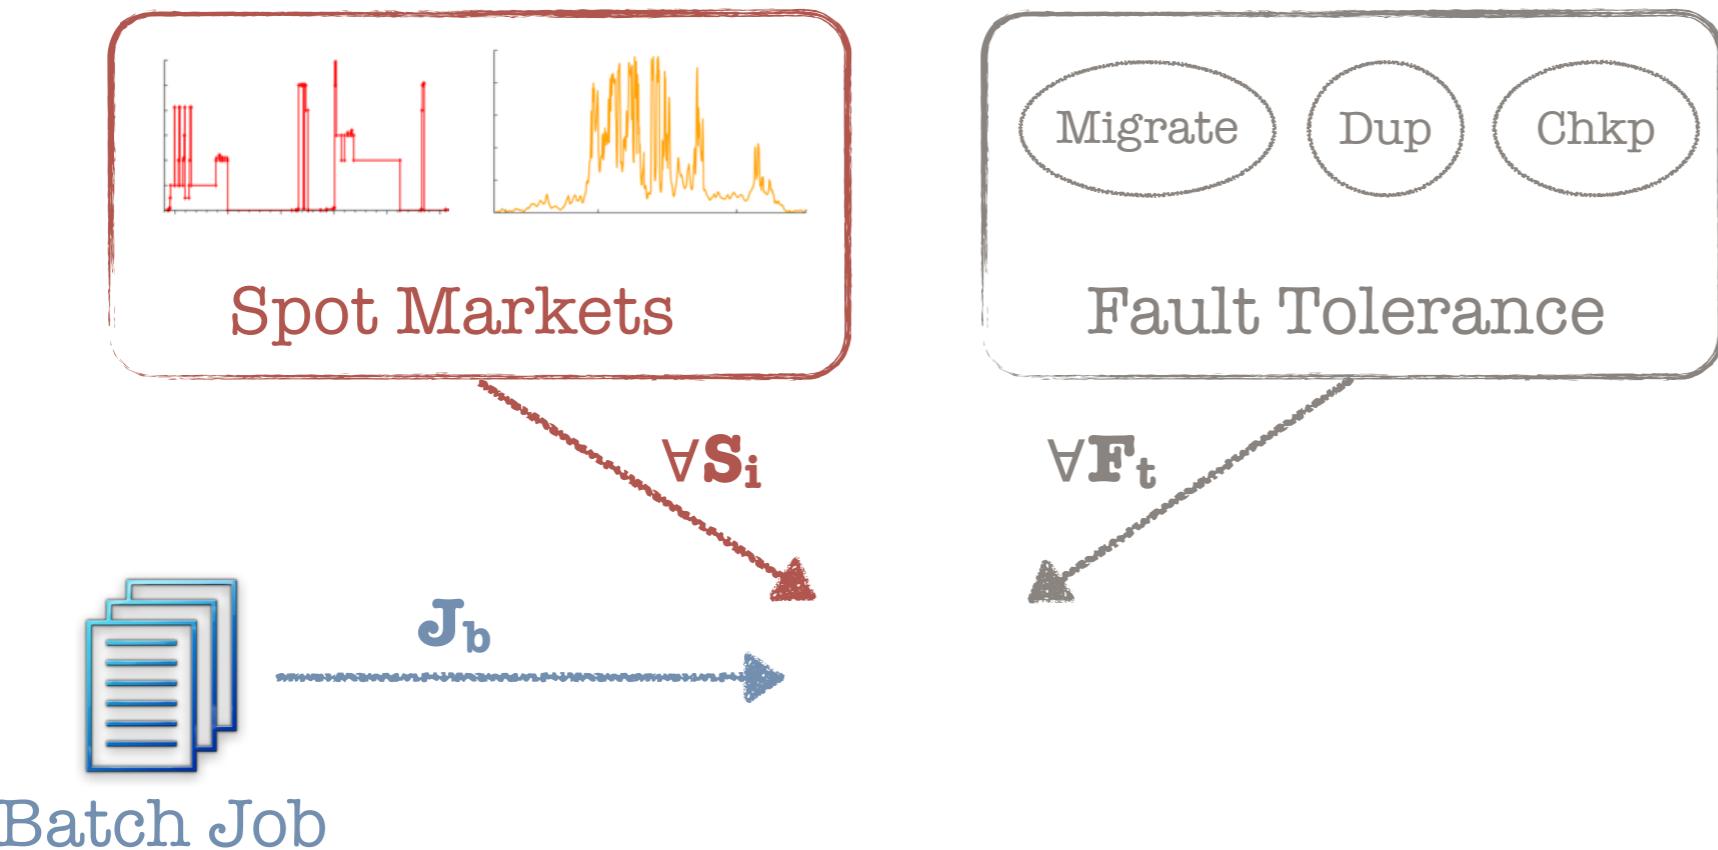
GREEDY SELECTION ALGORITHM

Selecting the best spot market and fault-tolerance mechanism



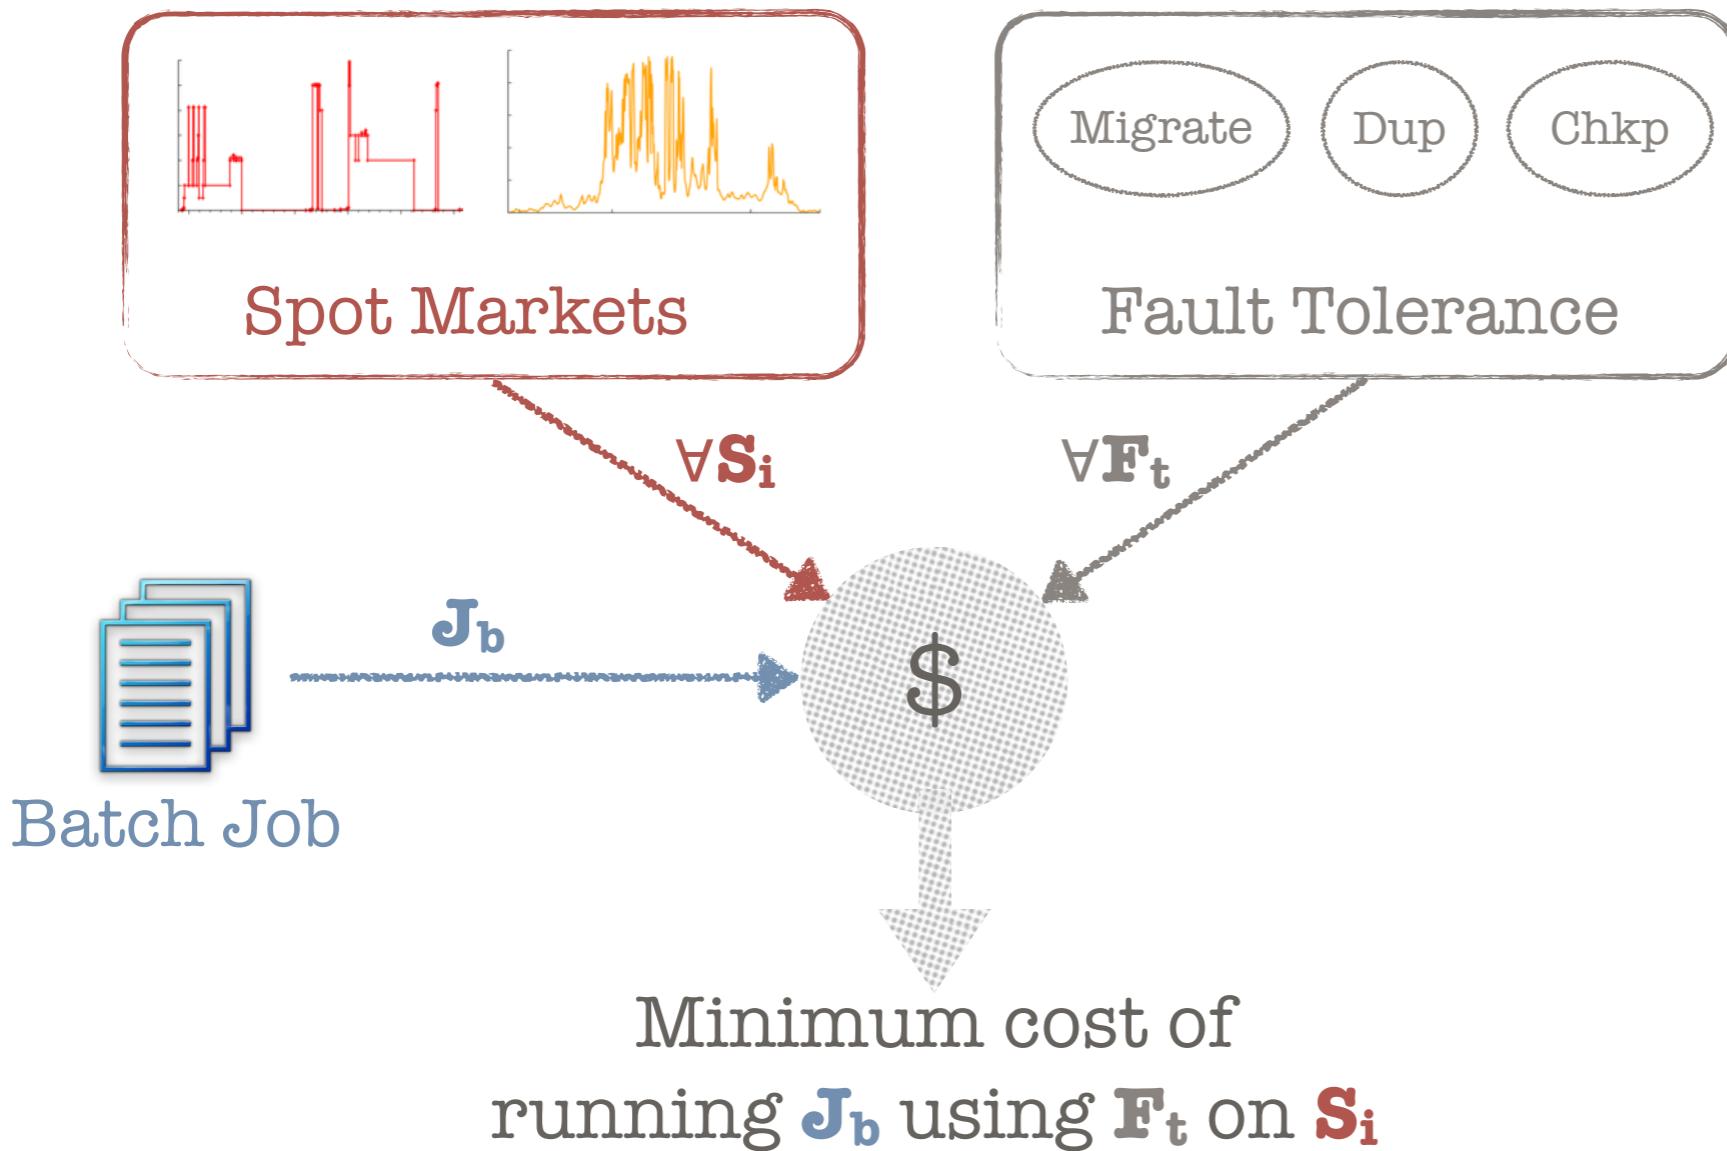
GREEDY SELECTION ALGORITHM

Selecting the best spot market and fault-tolerance mechanism



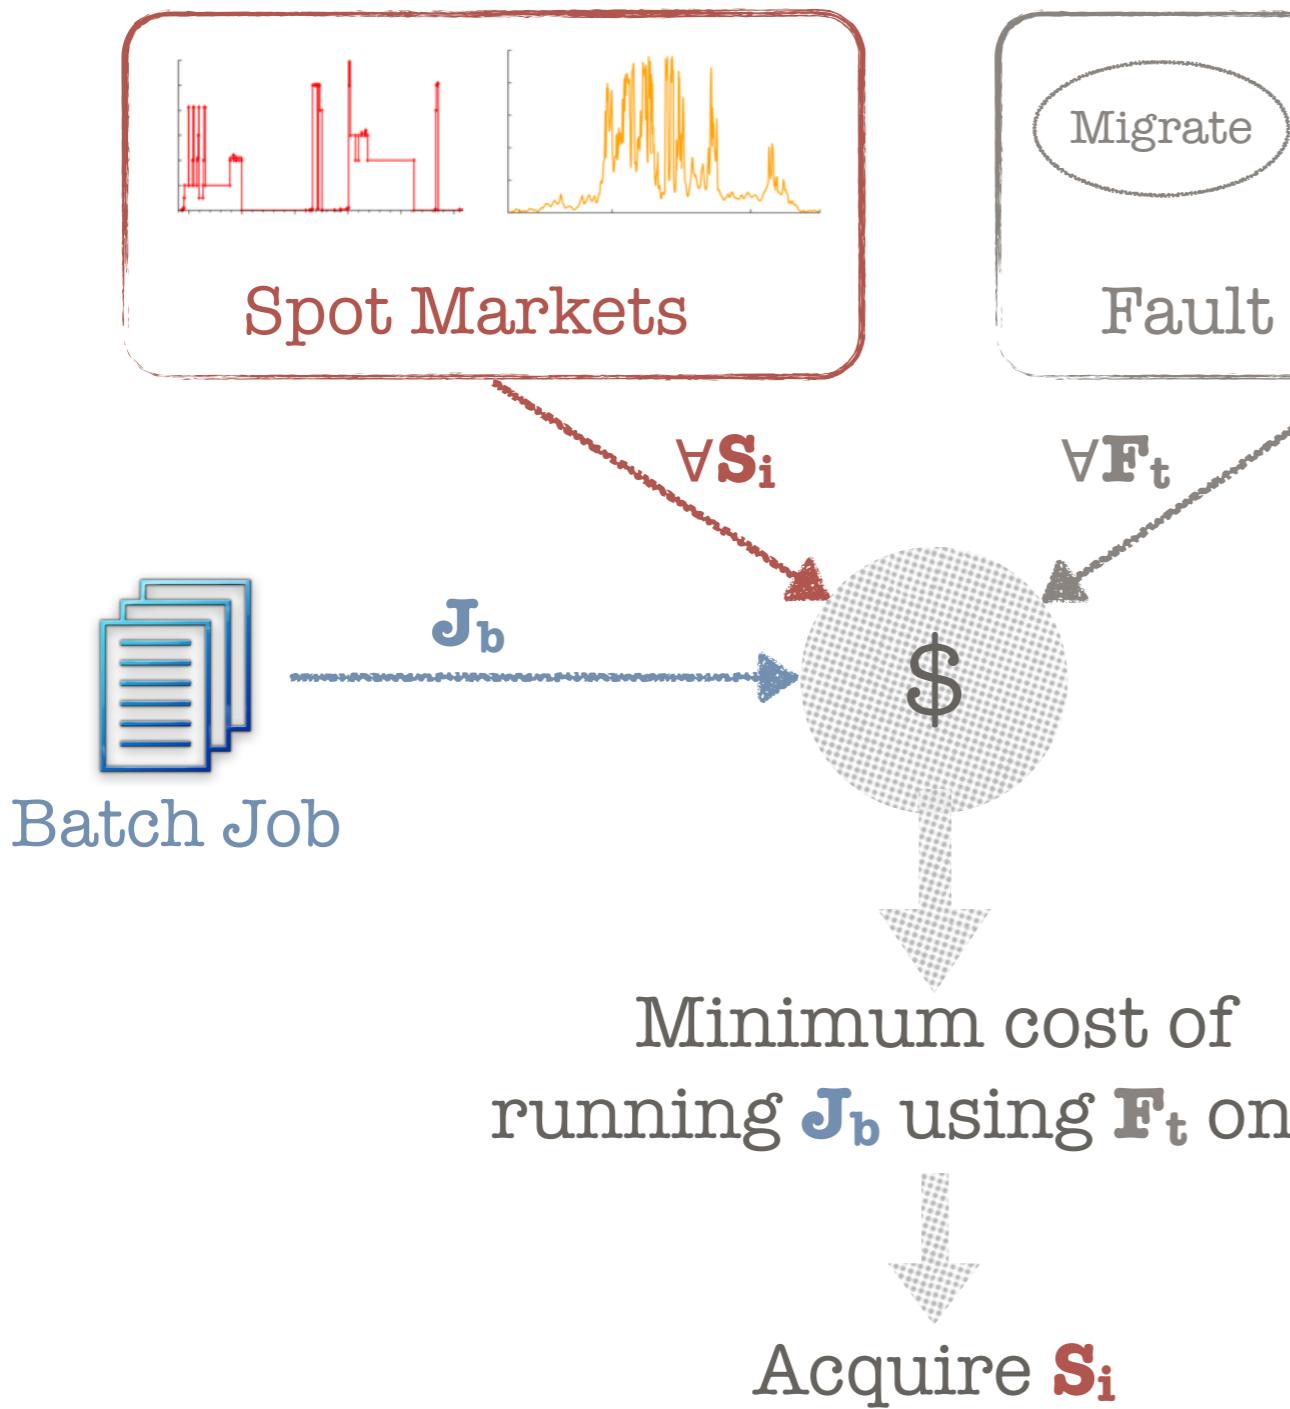
GREEDY SELECTION ALGORITHM

Selecting the best spot market and fault-tolerance mechanism



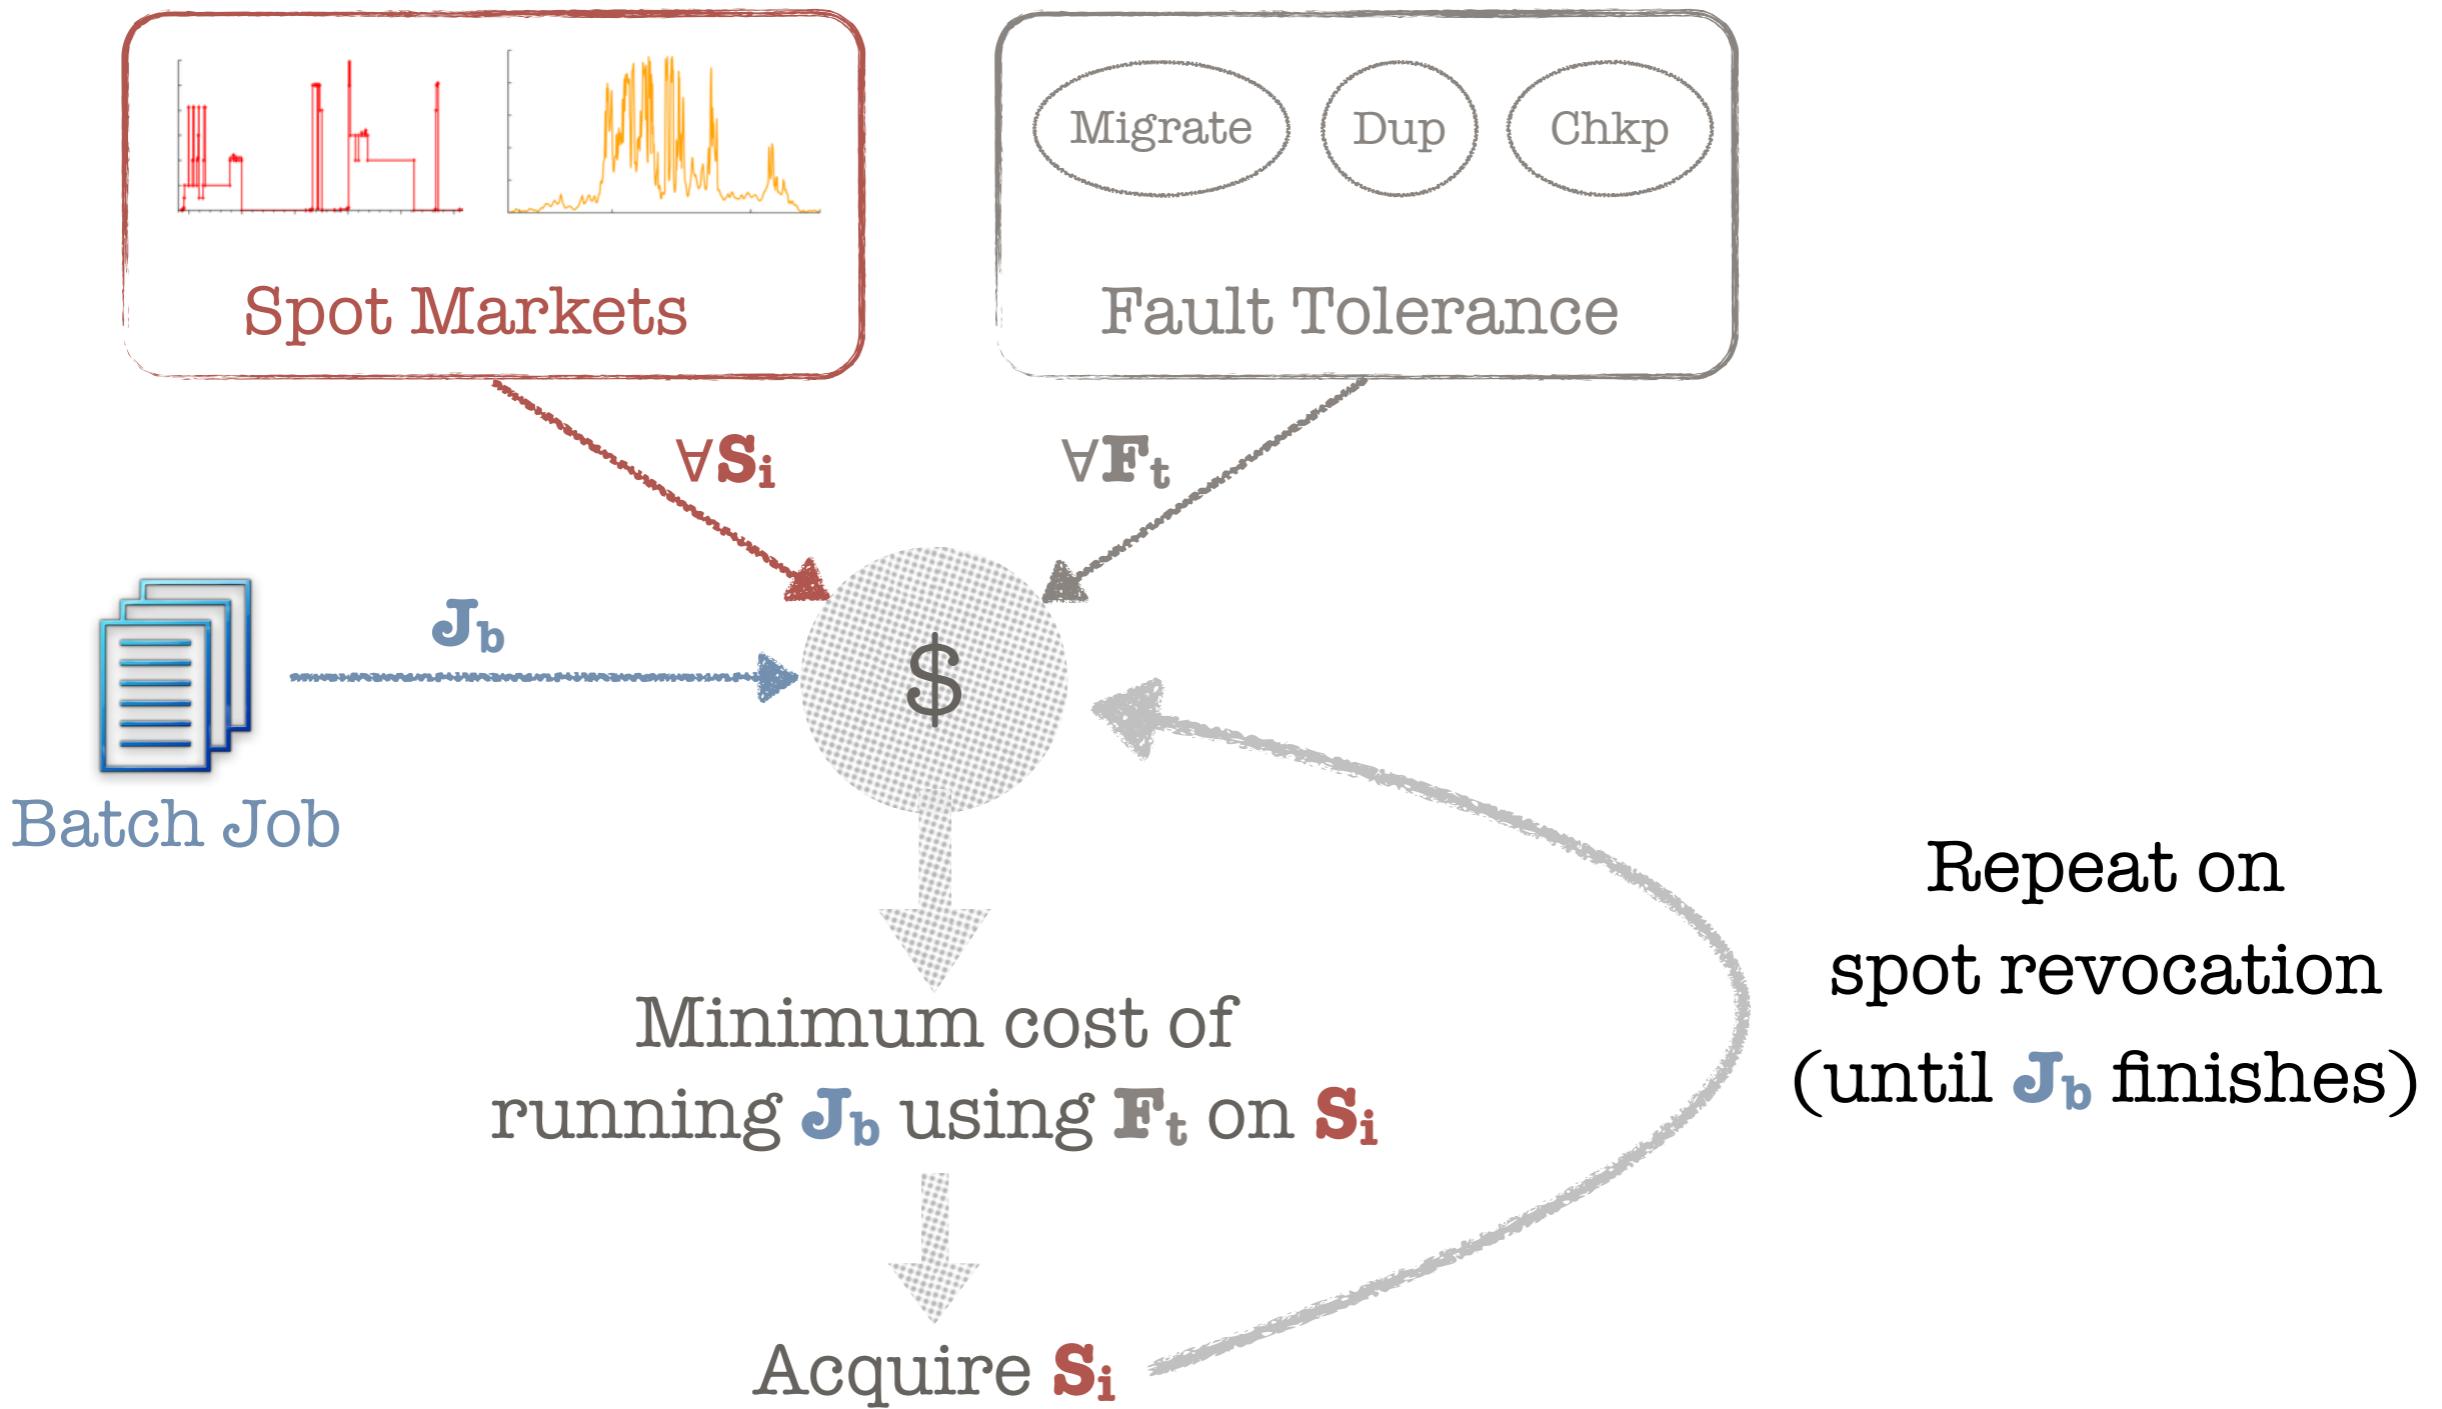
GREEDY SELECTION ALGORITHM

Selecting the best spot market and fault-tolerance mechanism



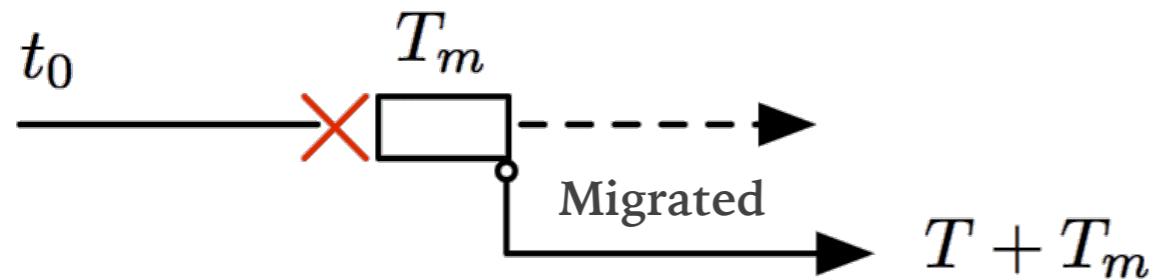
GREEDY SELECTION ALGORITHM

Selecting the best spot market and fault-tolerance mechanism



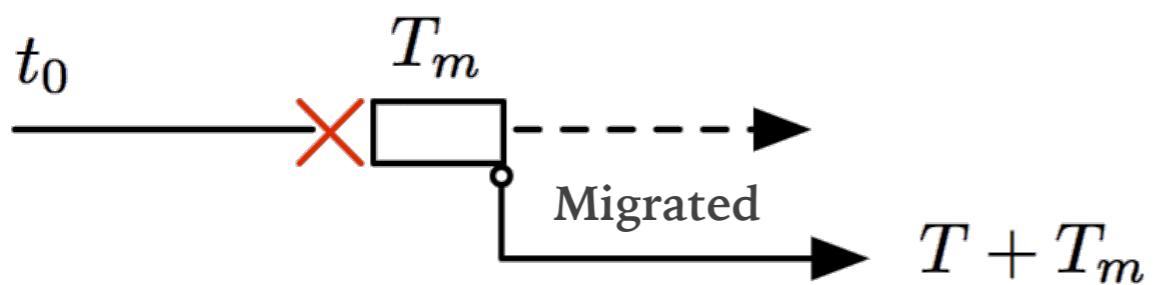
FAULT TOLERANCE (1/3)

Reactive Migration



FAULT TOLERANCE (1/3)

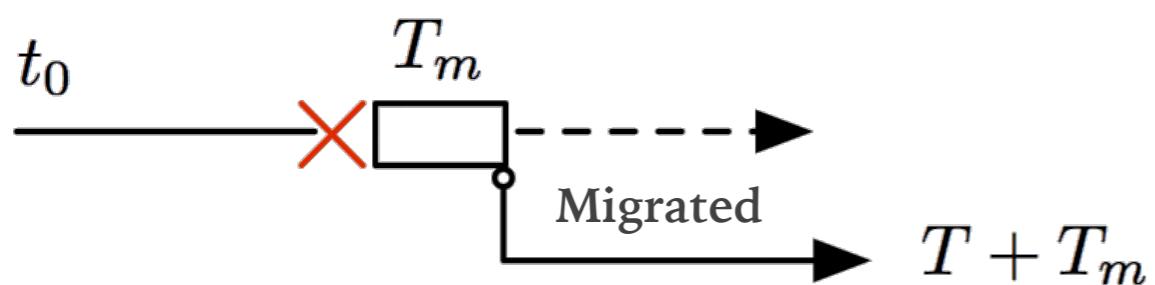
Reactive Migration



$$T_M \propto \frac{\text{size of memory + local disk}}{\text{remote disk bandwidth}}$$

FAULT TOLERANCE (1/3)

Reactive Migration



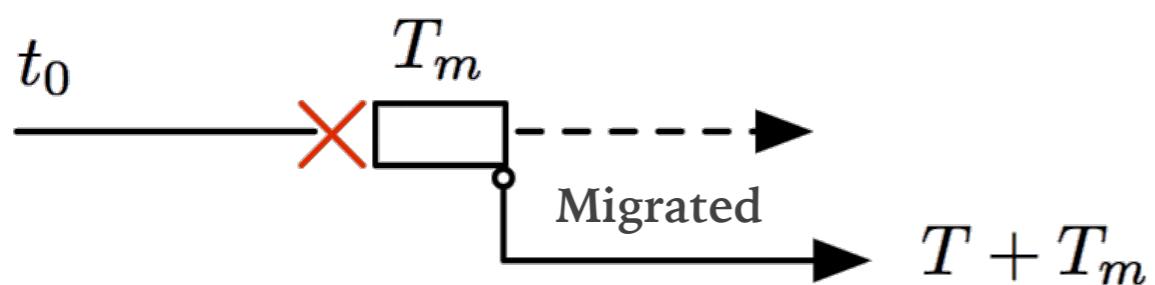
$$T_M \propto \frac{\text{size of memory + local disk}}{\text{remote disk bandwidth}}$$

$Z_k \rightarrow$ Random variable measuring time to revocation

$P_k \rightarrow$ Probability job gets revoked before completion

FAULT TOLERANCE (1/3)

Reactive Migration



$$T_M \propto \frac{\text{size of memory + local disk}}{\text{remote disk bandwidth}}$$

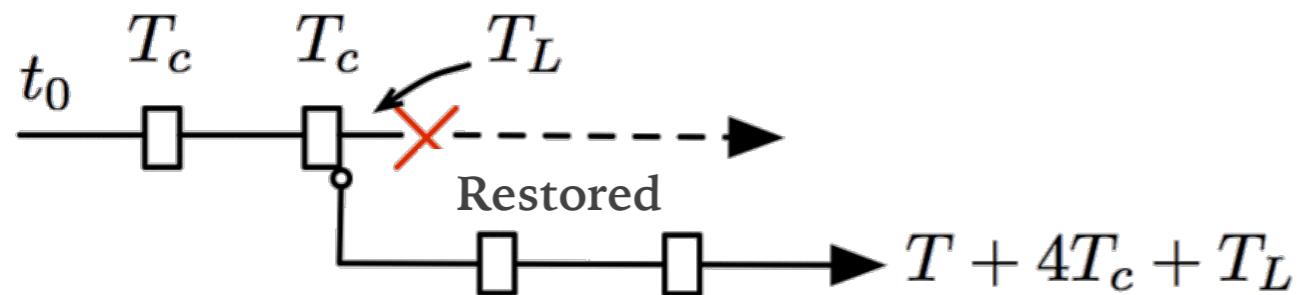
$Z_k \rightarrow$ Random variable measuring time to revocation

$P_k \rightarrow$ Probability job gets revoked before completion

$$\text{Cost} = \frac{E[\text{Price}_k]}{E[\text{Time}_k]} = \frac{[(1 - P_k) * T + P_k * (E(Z_k) + T_M)] * \text{spot-price}}{(1 - P_k) * T + P_k * E(Z_k)}$$

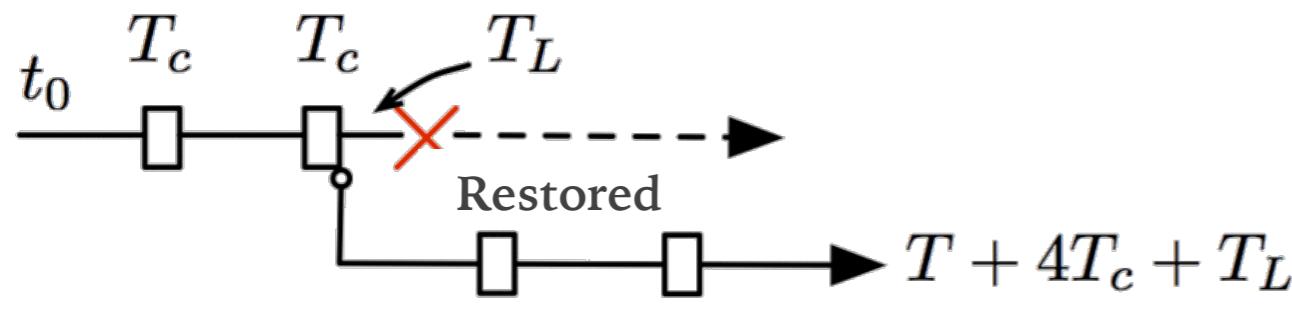
FAULT TOLERANCE (2/3)

Proactive Checkpoint



FAULT TOLERANCE (2/3)

Proactive Checkpoint



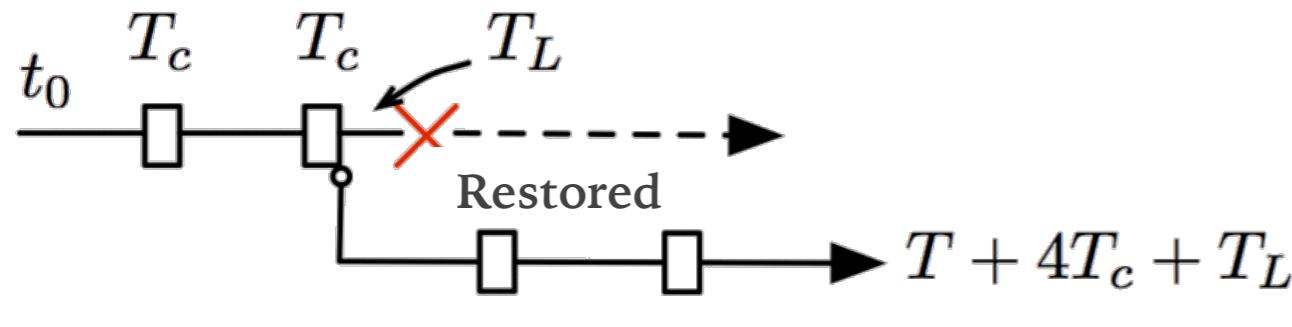
$$T_C \propto \frac{\text{size of memory} + \text{local disk}}{\text{remote disk bandwidth}}$$

$$T_L \propto (\tau/2)$$

τ : checkpoint frequency

FAULT TOLERANCE (2/3)

Proactive Checkpoint



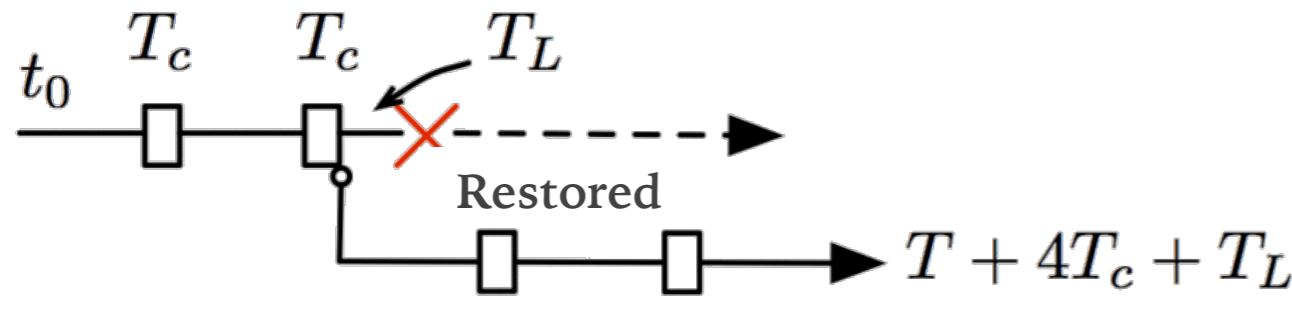
$$T_C \propto \frac{\text{size of memory} + \text{local disk}}{\text{remote disk bandwidth}}$$
$$T_L \propto (\tau/2)$$

τ : checkpoint frequency

$$\text{Total Overhead} = (T/\tau) * T_C + T_L$$

FAULT TOLERANCE (2/3)

Proactive Checkpoint



$$T_C \propto \frac{\text{size of memory} + \text{local disk}}{\text{remote disk bandwidth}}$$
$$T_L \propto (\tau/2)$$

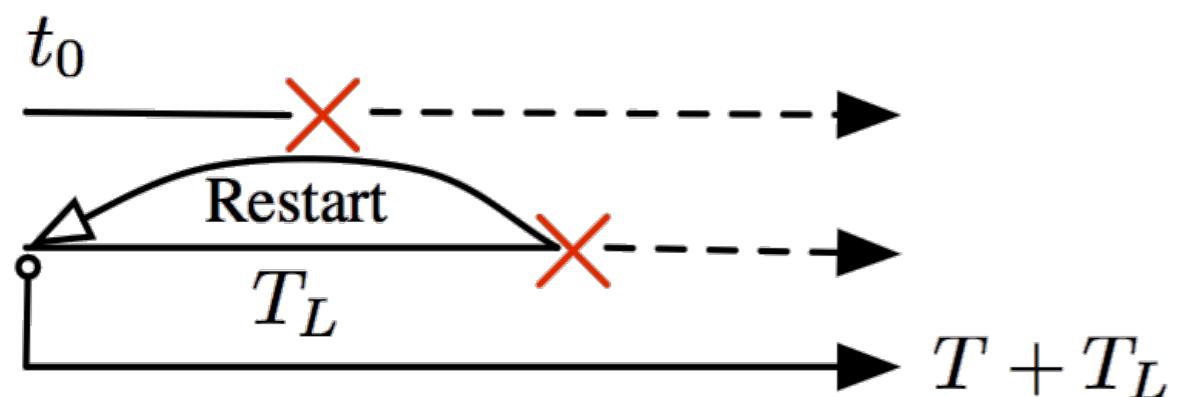
τ : checkpoint frequency

$$\text{Total Overhead} = (T/\tau) * T_C + T_L$$

Cost overhead is primarily a function of job's resource usage

FAULT TOLERANCE (3/3)

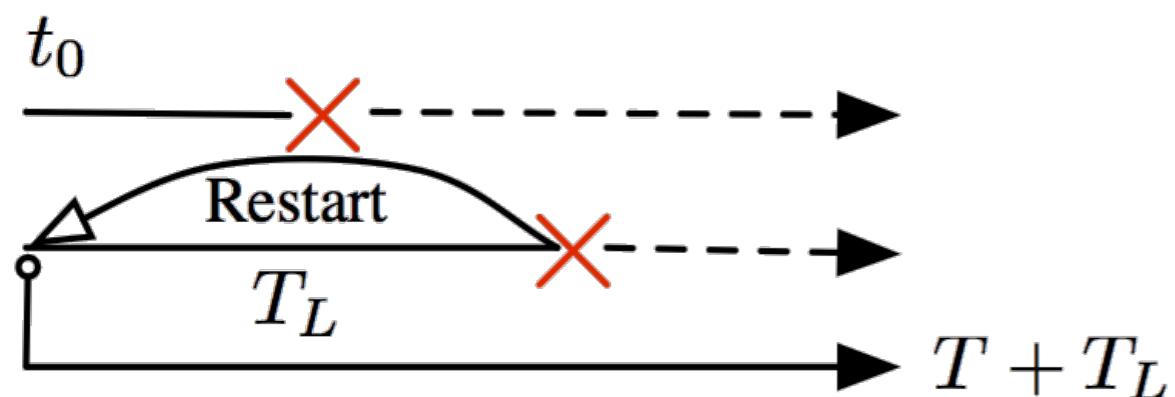
Spot Replication



$T_L \propto$ market volatility

FAULT TOLERANCE (3/3)

Spot Replication



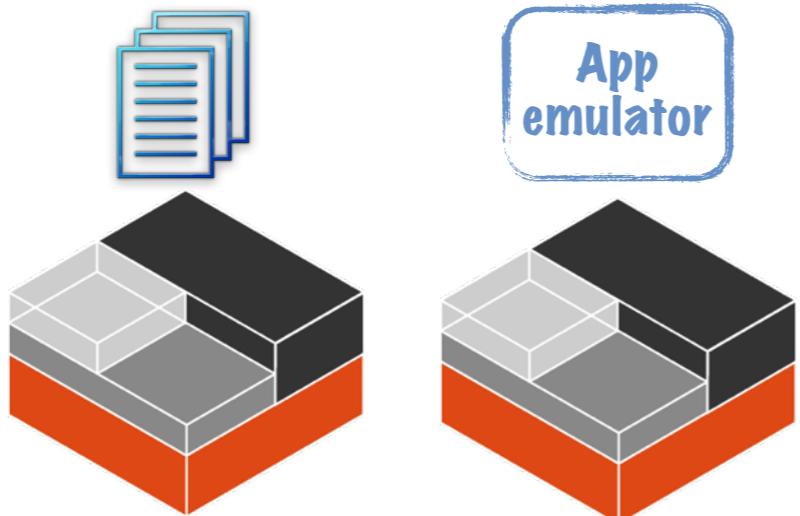
$T_L \propto$ market volatility

Total Overhead \propto (Replication factor, T_L)

Cost overhead is primarily a function of market characteristics

EVALUATION (1/3)

SpotOn Prototype



- Built on *Linux Containers* for efficient checkpointing / migration
- *App Emulator* to create synthetic jobs with varying resource usage

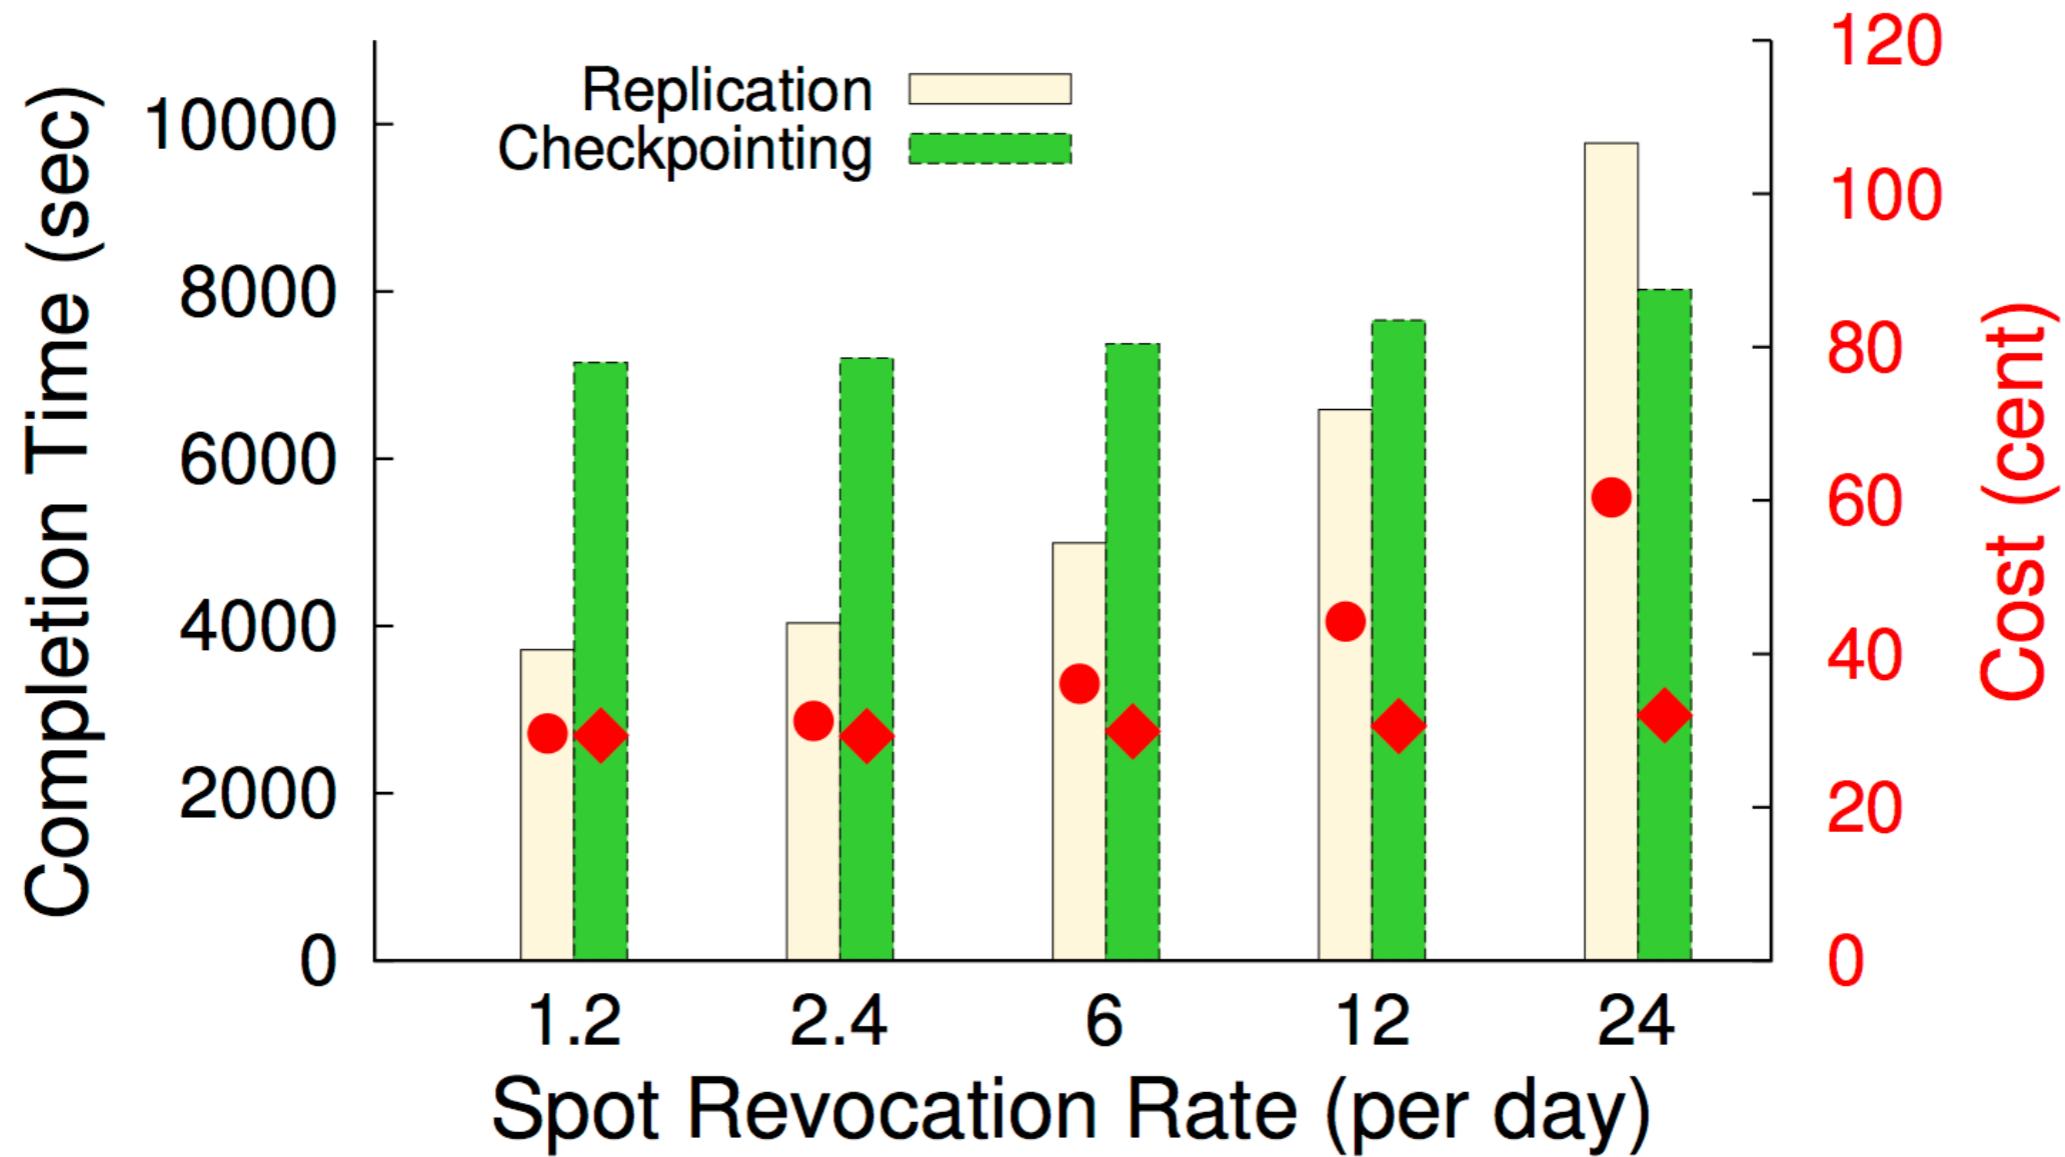


EVALUATION (2/3)

Effect of Application and Spot Market Characteristics

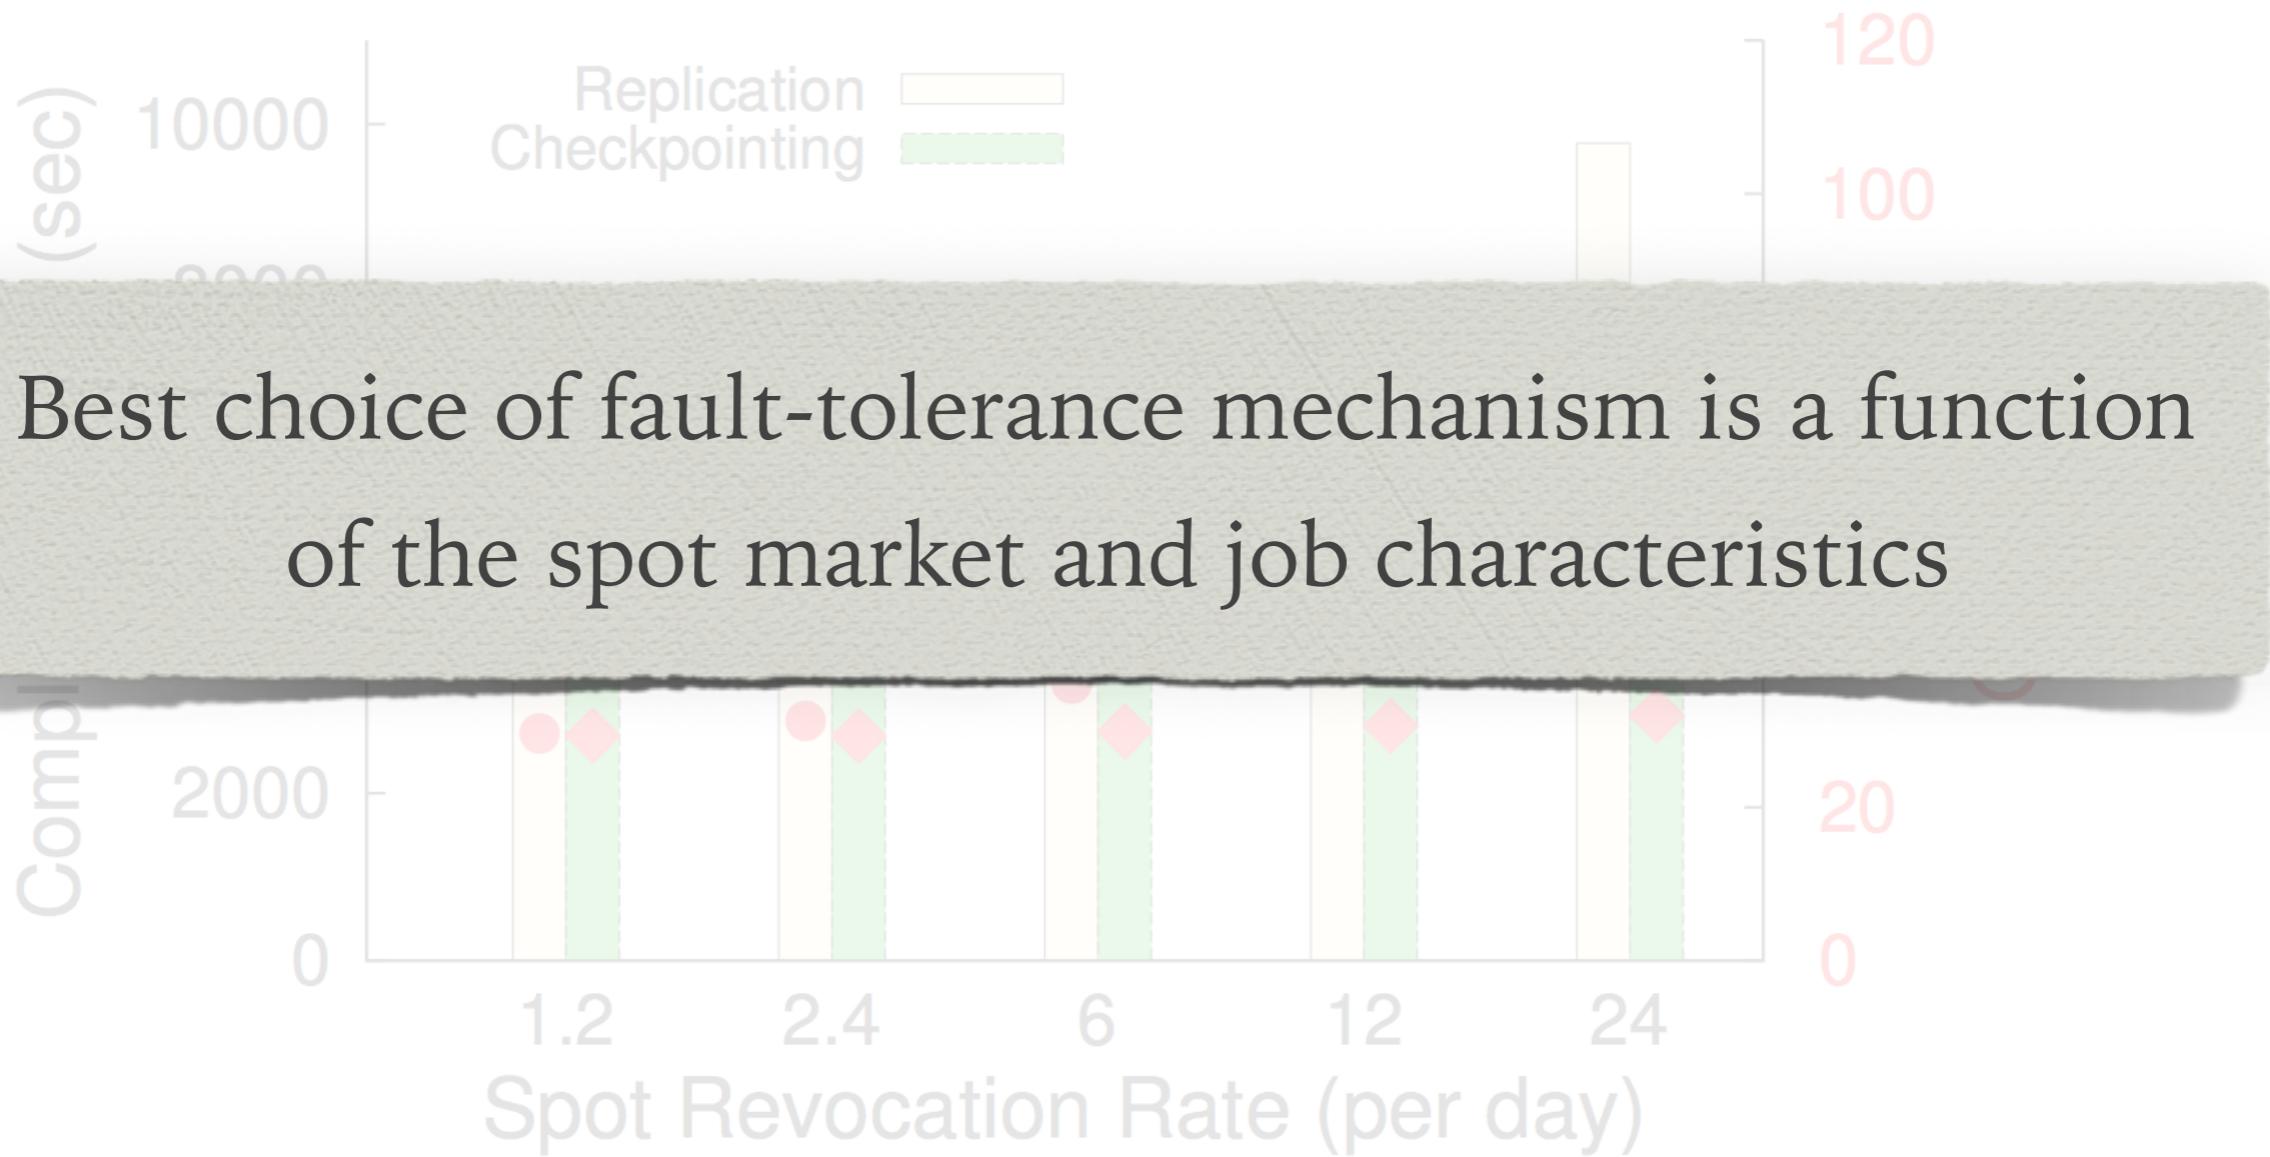
EVALUATION (2/3)

Effect of Application and Spot Market Characteristics



EVALUATION (2/3)

Effect of Application and Spot Market Characteristics

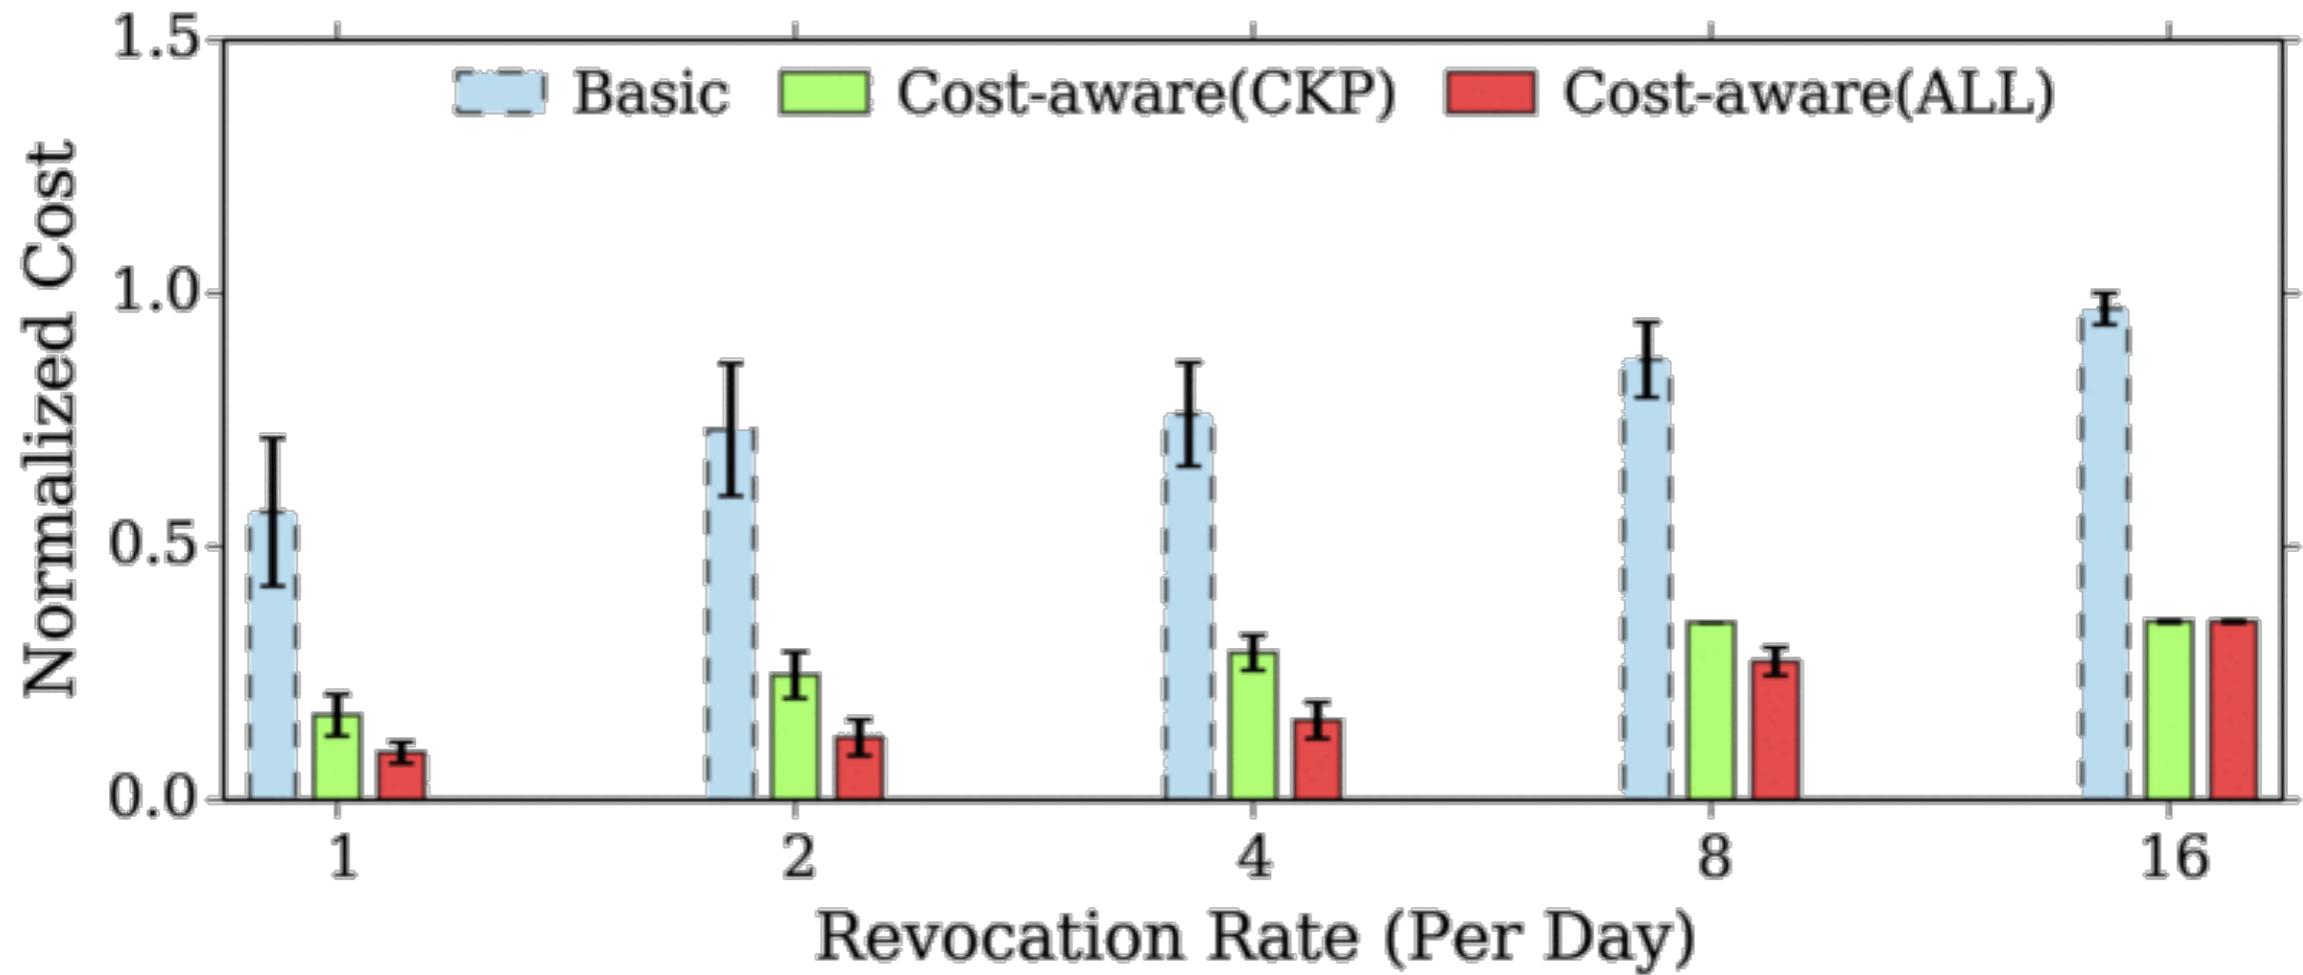


EVALUATION (3/3)

Effect of Cost-aware Selection

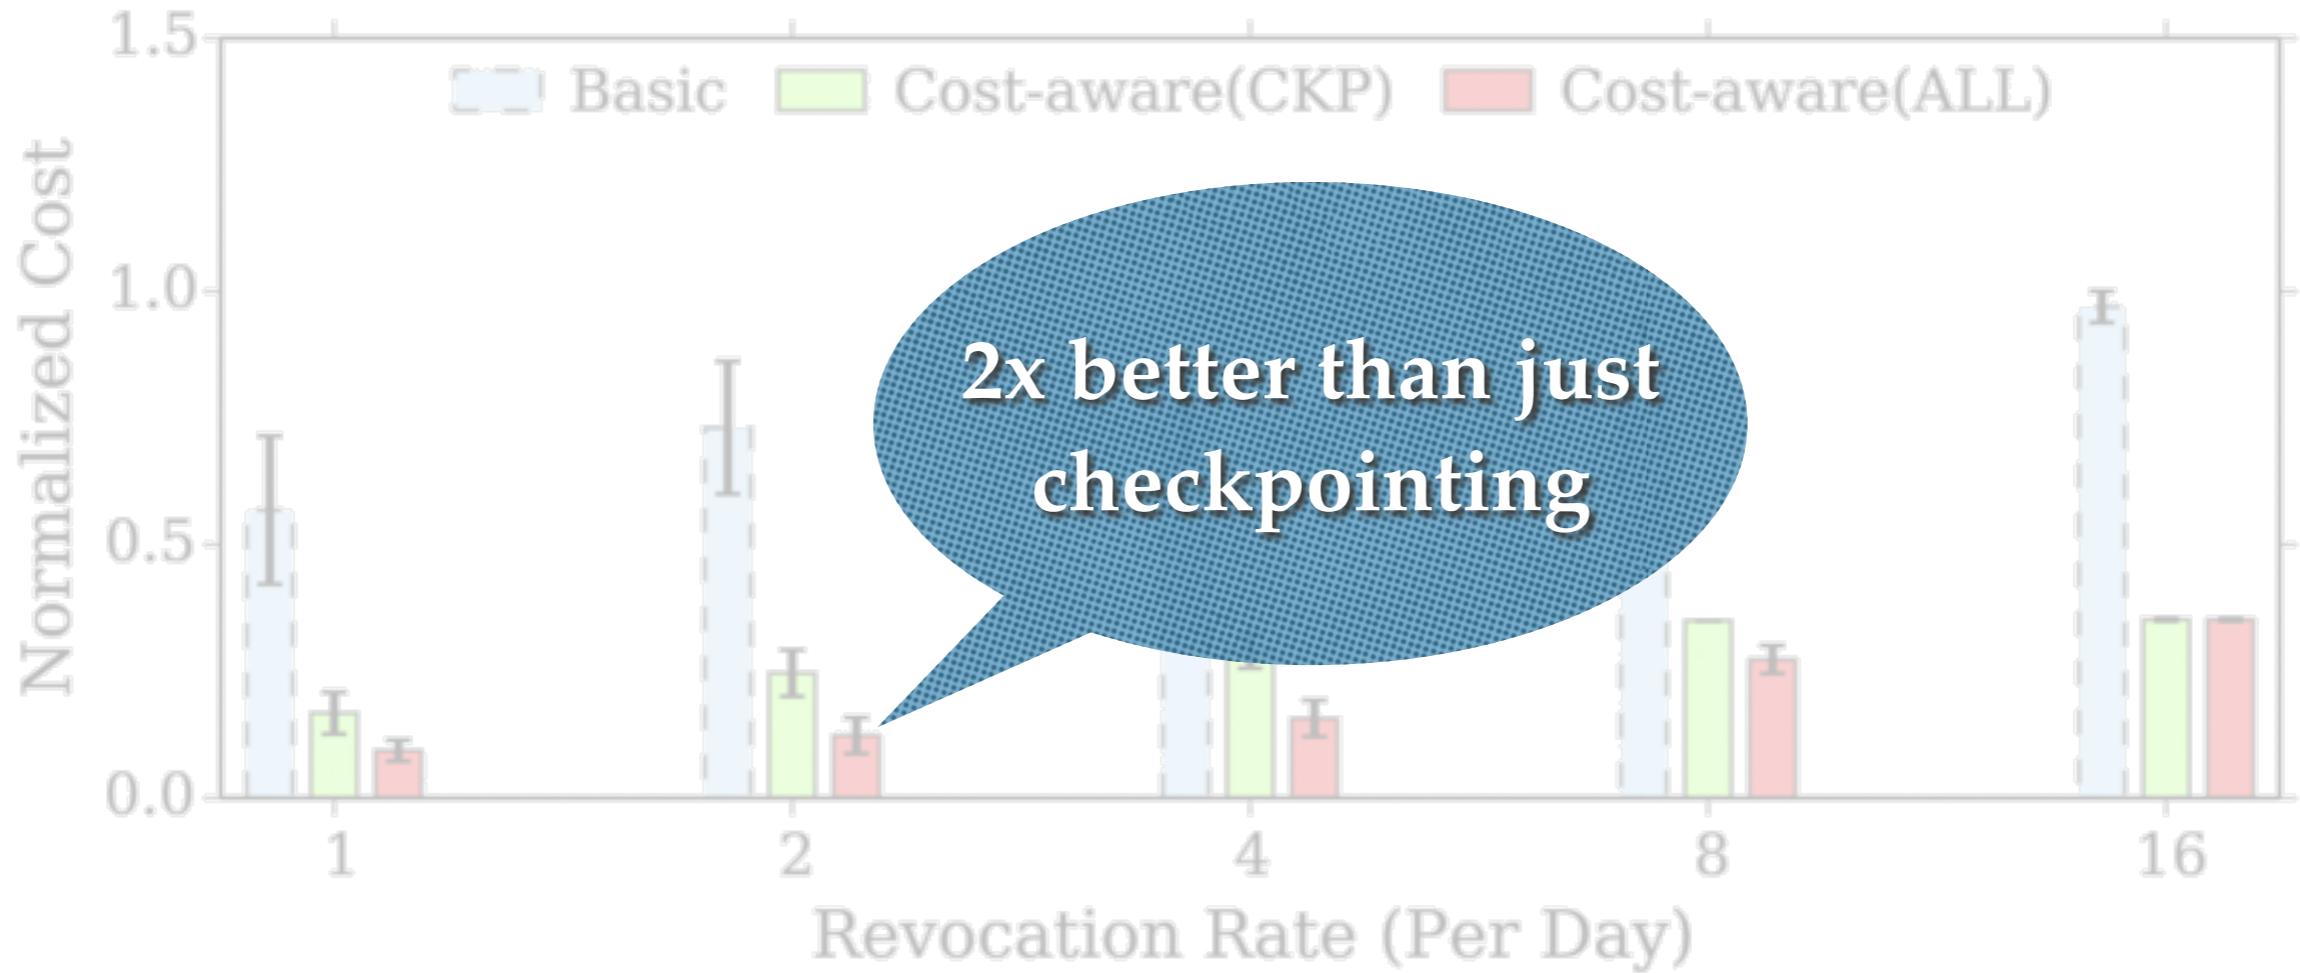
EVALUATION (3/3)

Effect of Cost-aware Selection



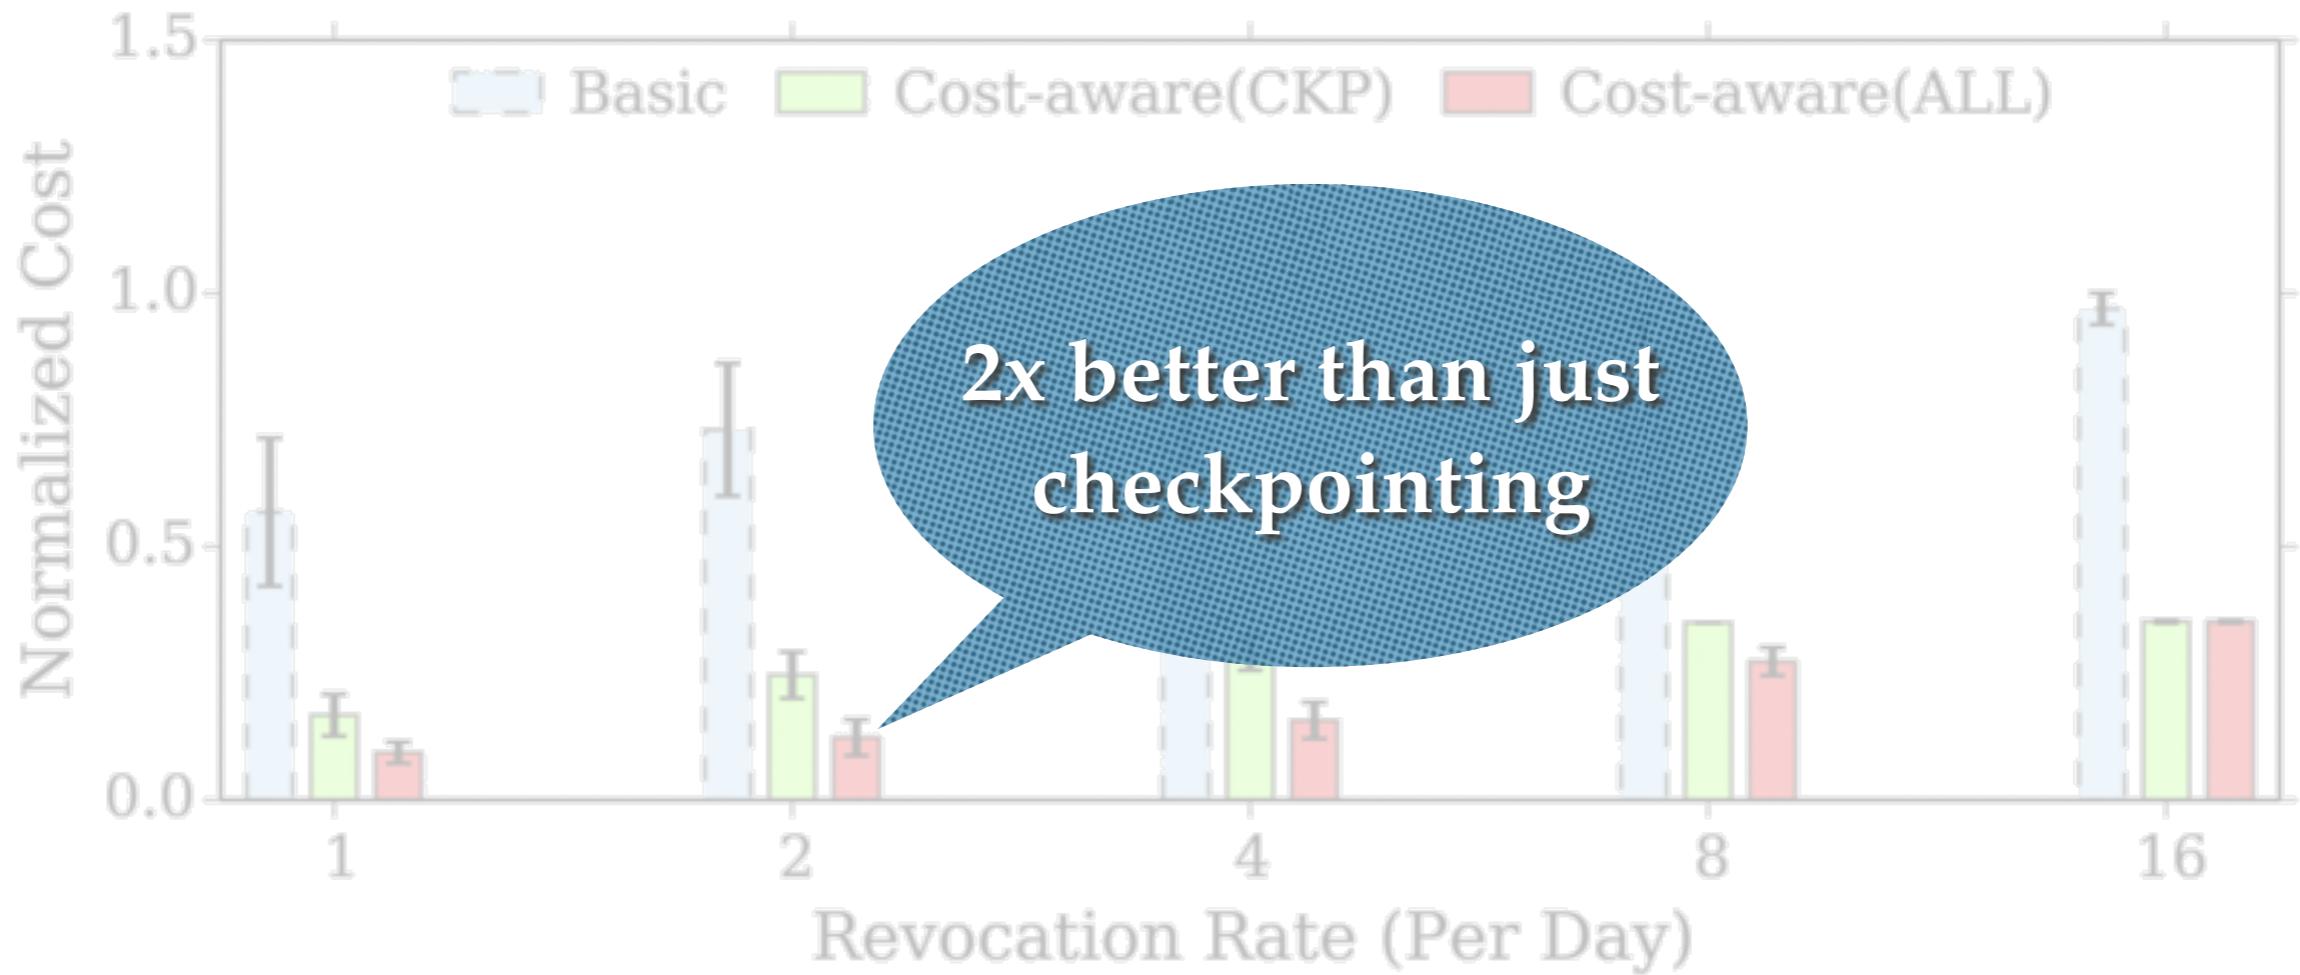
EVALUATION (3/3)

Effect of Cost-aware Selection



EVALUATION (3/3)

Effect of Cost-aware Selection



On Google cluster trace, Cost-aware selection achieved
91.9% savings with little impact on performance

CONCLUSION

- Spot markets offer **arbitrage** opportunities
- SpotOn manages application and market complexities
- We model **fault tolerance** and propose a **selection algorithm**
- Prototype** on Amazon EC2
- Achieves ~90% cost savings

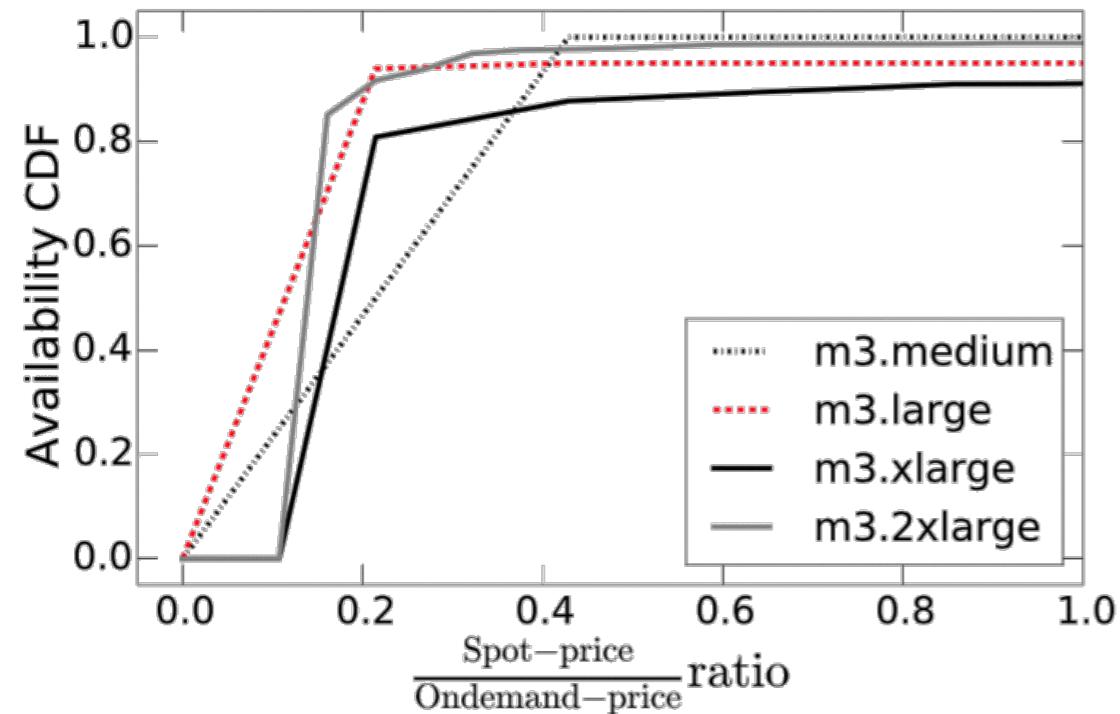
CONCLUSION

- Spot markets offer **arbitrage** opportunities
- SpotOn manages application and market complexities
- We model **fault tolerance** and propose a **selection algorithm**
- Prototype** on Amazon EC2
- Achieves ~90% cost savings

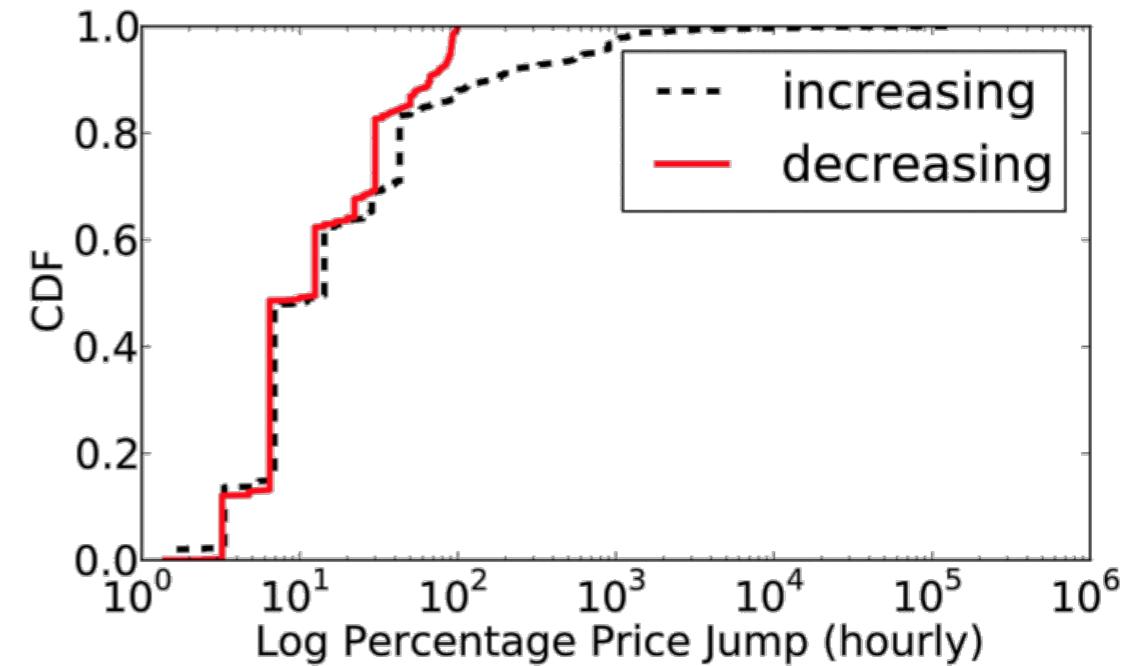
Thank you!

BACKUP SLIDES

BIDDING

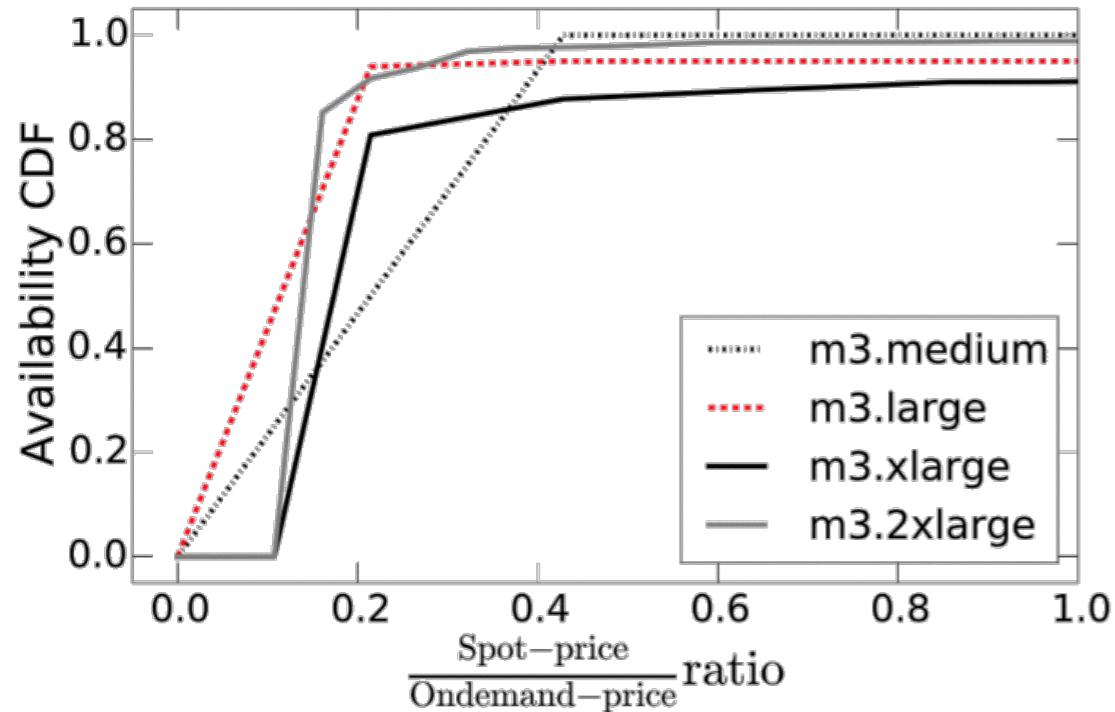


Spot price distribution has long tail

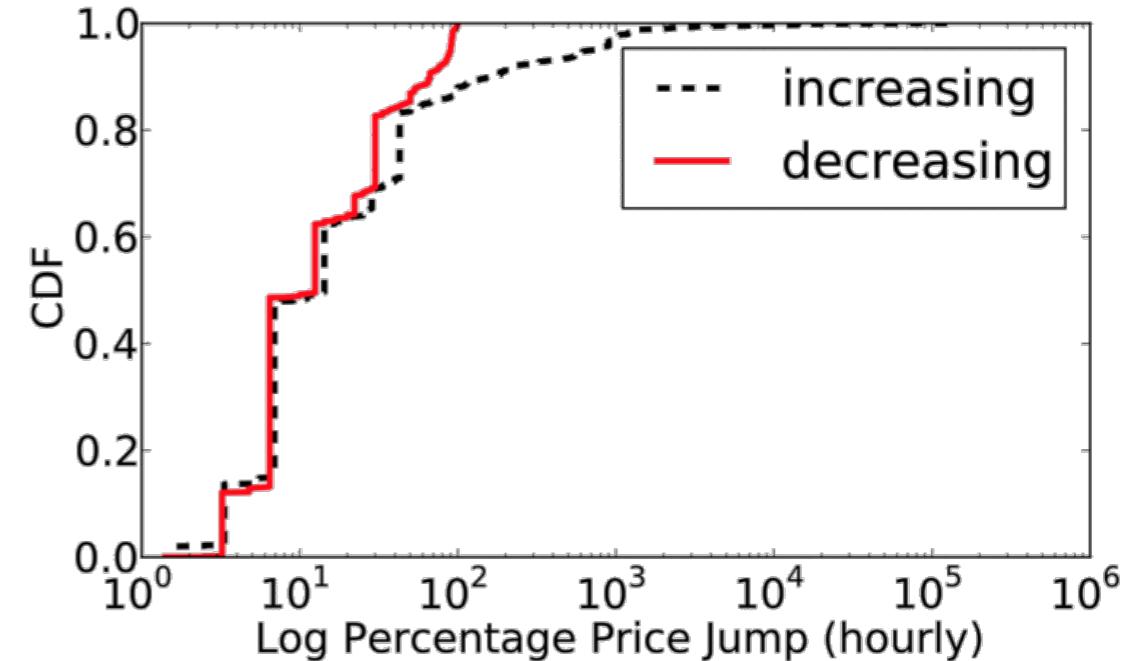


Spot price changes are peaky

BIDDING



Spot price distribution has long tail



Spot price changes are peaky

- 📌 Bidding does affect volatility but *not drastically*
- 📌 SpotOn always bids at the on-demand price
If spot price goes above on-demand, SpotOn would choose on-demand